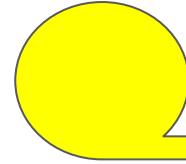


Using and extending the Kubernetes API programmatically with Go

A workshop at GopherCon UK 2018

Stefan Schimanski sttts@redhat.com

Michael Hausenblas mhausenb@redhat.com



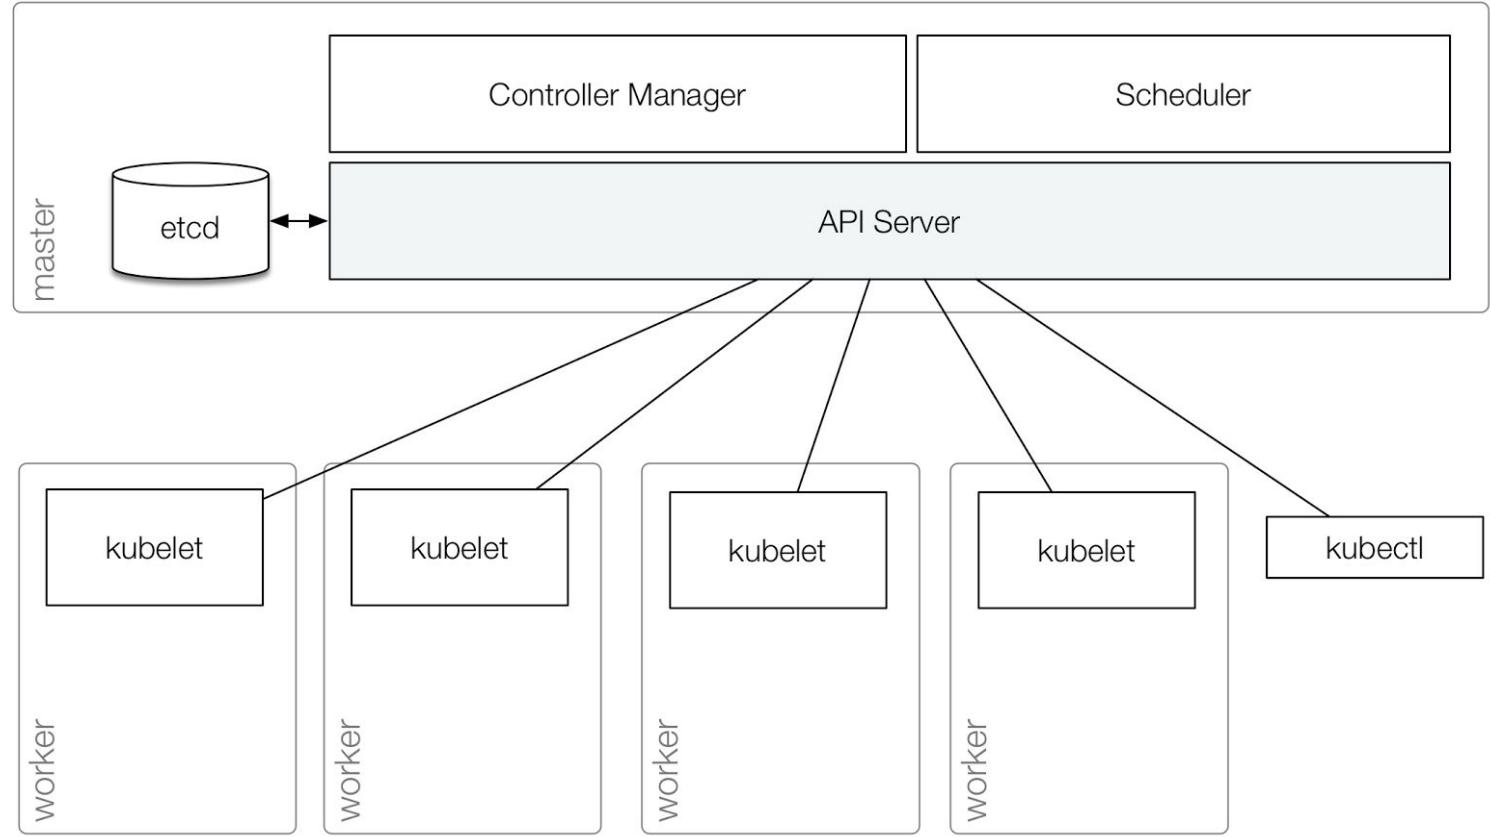
Slides

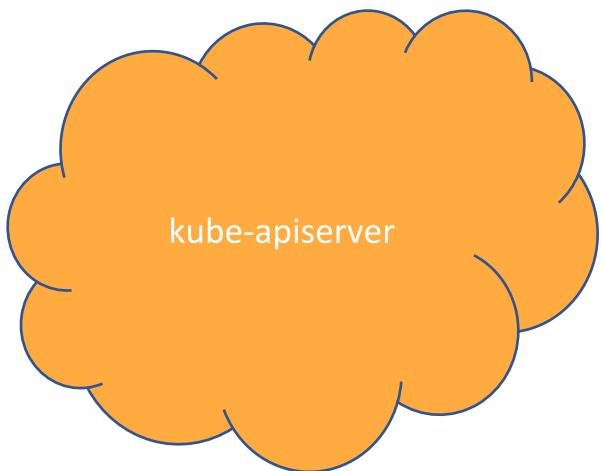
<https://301.sh/2018-gopherconuk-slides>

Agenda

- **API basics:** apigroups, resource vs. kind, versioning, discovery, **curl scenario**
- **Client-go basics:** clientsets, kubeconfig, **client-go scenario**, informers, **informer scenario**
- **Kubernetes objects in Golang:** ObjectMeta, TypeMeta
- **API Machinery:** Scheme, GVK, GVR, RestMapper
- **CRDs from user point of view:** CRD yaml, Validation, **CRD yaml scenario**
- **Code-Generation:** deepcopy-gen, client-gen, informer-gen, lister-gen
- **CRDs from developer point of view:** types.go, register.go, pkg/apis, code-generator, **CRD Go scenario**
- **Controllers:** controller loop, optimistic concurrency, edge-vs-level driven, **controller scenario**
- **Kubebuilder:** init, create api, **kubebuilder scenario – trainer led**
- **Operator SDK:** new operator, **operator scenario – trainer led**
- **Other topics:** deployment / in-cluster config, RBAC rules, service account, subresources, versioning, **misc scenario ("deploy your controller in the cluster")**

Kubernetes API basics





Restful http API

/
/version

/api
/api/v1/**pods**
/api/v1/**pods**/**<name>**
/api/v1/**pods**/**<name>**/status

/apis
/apis/batch
/apis/batch/v2alpha1
/apis/batch/v2alpha1/**jobs**
/apis/batch/v2alpha1/**cronjobs**
/apis/batch/v1
/apis/batch/v1/**jobs**

User:

```
$ kubectl create -f  
foo.yaml
```

Node:

proxy

kubelet

Pods:

cloud native apps

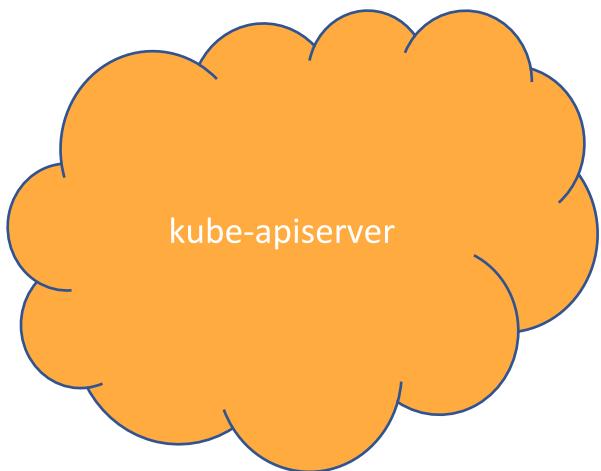
Master:

scheduler

controller
manager

ingress
controller

apiserver



Restful http API

/
/version

/api
/api/v1/**pods**
/api/v1/**pods**/**<name>**
/api/v1/**pods**/**<name>**/status

/apis
/apis/batch
/apis/batch/v2alpha1
/apis/batch/v2alpha1/**jobs**
/apis/batch/v2alpha1/**cronjobs**
/apis/batch/v1
/apis/batch/v1/**jobs**

/apis/<our-group>/v1/<our-resource>

User:

```
$ kubectl create -f  
foo.yaml
```

Node:

proxy

kubelet

Pods:

cloud native apps

Master:

scheduler

controller
manager

ingress
controller

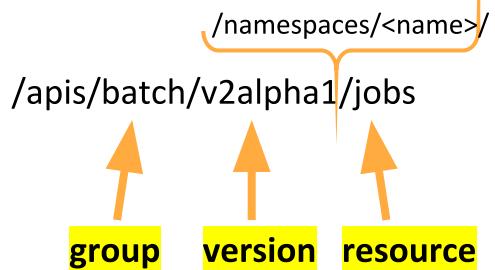
apiserver

```
$ kubectl proxy
$ curl http://127.0.0.1:8001/apis
$ kubectl get --raw=/apis

$ kubectl annotate namespace default workshop=gopherconuk
$ kubectl get namespace default -o json
$ curl http://127.0.0.1:8001/api/v1/namespaces/default
{
  "kind": "Namespace",
  "apiVersion": "v1",
  "metadata": {
    "name": "default",
    "selfLink": "/api/v1/namespaces/default",
    "uid": "f86e4e1f-94ad-11e8-9d48-0242ac11002b",
    "resourceVersion": "283",
    "creationTimestamp": "2018-07-31T10:39:15Z",
    "annotations": {
      "workshop": "gopherconuk"
    }
  },
  "spec": {
    "finalizers": [
      "kubernetes"
    ]
  },
  "status": {
    "phase": "Active"
  }
}
```

```
$ curl http://127.0.0.1:8001/api/v1/namespaces/default | \
jq ".metadata.annotations[\"workshop\"] = \"$(date)\"" | \
curl -H "Content-Type: application/json" -X PUT -d @- http://127.0.0.1:8001/api/v1/namespaces/default
% Total % Received % Xferd Average Speed Time Time Time Current
          Dload Upload Total Spent Left Speed
100 430 100 430 0 0 29689 0 --:--:--:--:--:--:--:-- 30714
{
  "kind": "Namespace",
  "apiVersion": "v1",
  "metadata": {
    "name": "default",
    "selfLink": "/api/v1/namespaces/default",
    "uid": "69c4fd62-94a9-11e8-b75a-0242ac11006e",
    "resourceVersion": "2091",
    "creationTimestamp": "2018-07-31T10:06:37Z",
    "annotations": {
      "workshop": "Tue Jul 31 10:32:16 UTC 2018"
    }
  },
  "spec": {
    "finalizers": [
      "kubernetes"
    ]
  },
  "status": {
    "phase": "Active"
  }
}
```

HTTP paths:



Possible (logical) verbs:

- GET
- CREATE (= http POST)
- UPDATE (= http PUT)
- LIST (http GET on root path)
- DELETE
- PATCH
- WATCH (?watch=true)

```
{  
  "apiVersion": "v2alpha1",  
  "kind": "Job",  
  "metadata": {  
    "name": "backup"  
  },  
  "spec": { ... }  
}
```

/apis/batch/v2alpha1/jobs

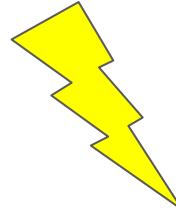
Resource vs. Kind

http path vs. logical object

```
{  
  "apiVersion": "v1",  
  "kind": "Status",  
}
```

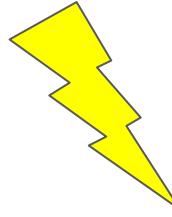
```
{  
  "apiVersion": "v2alpha1",  
  "kind": "Job",  
  "metadata": {  
    "name": "backup"  
  },  
  "spec": { ... }  
}
```

* I omitted the namespace in /apis/batch/v1/jobs/**namespaces/default**/backup



Katacoda

<https://301.sh/2018-gopherconuk-katacoda>



Katacoda

→ API basics (step 2)

Client-go basics

k8s.io/client-go

kubernetes / client-go

Code Issues 40 Pull requests 0 Projects 0 Wiki Insights

Go client for Kubernetes.

1,242 commits 11 branches 137 releases 178 contributors Apache-2.0

Branch: master New pull request Create new file Upload files Find file Clone or download

k8s-publishing-bot Merge pull request #66518 from deads2k/dynamic-02-ordie ... Latest commit 3db81bd 2 days ago

.github client-go: document README exception in .github/PULL_REQUEST_TEMPLATE.md 2 months ago

Godeps Merge pull request #66518 from deads2k/dynamic-02-ordie 2 days ago

deprecated-dynamic fix dynamic client name 3 months ago

discovery fix info level message 6 months ago

dynamic add missing OrDie variant for dynamic client construction 3 days ago

examples client-go/examples/fake-client: add doc.go to fix go build warnings a month ago

informers Update generated files 27 days ago

kubernetes Run code gen 23 days ago

listers Autogenerated stuff 2 months ago

pkg generated: Avoid use of reflect.Call in conversion code paths 23 days ago

plugin/pkg/client/auth Add missing error handling in schema-related code 2 months ago

rest client-go: fix error message spelling in rest config 10 days ago

restmapper fill in normal restmapping info with the legacy guess 10 days ago

scale generated: Avoid use of reflect.Call in conversion code paths 23 days ago

testing add Patch support in fake kubeClient 4 months ago

third_party/forked/golang/template sync: initially remove files BUILD */BUILD BUILD.bazel */BUILD.bazel 4 months ago

tools client-go: update documentation for remotecommand.StreamOptions 10 days ago

transport Add TLS support to exec authenticator plugin 2 months ago

util apiserver: use fixtures for self-signed certs in test server 20 days ago

clients →

clientcmd →

[Code](#)[Issues 0](#)[Pull requests 0](#)[Projects 0](#)[Wiki](#)[Insights](#)

The canonical location of the Kubernetes API definition.

4,764 commits

5 branches

100 releases

307 contributors

Apache-2.0

Branch: master ▾

[New pull request](#)[Create new file](#)[Upload files](#)[Find file](#)[Clone or download ▾](#)

 k8s-publishing-bot	Merge pull request #66047 from krunaljain/bugfix/csi_default_fs_type	Latest commit 183f332 16 days ago
 .github	Treat staging repos as authoritative for all files	7 months ago
 Godeps	Merge pull request #65499 from krunaljain/bugfix/csi_default_fs_type	16 days ago
 admission/v1beta1	Update generated files	a month ago
 admissionregistration	Update generated files	a month ago
 apps	Update generated files	27 days ago
 authentication	Update generated files	a month ago
 authorization	Update generated files	a month ago
 autoscaling	Update generated files	27 days ago
 batch	Update generated files	27 days ago
 certificates	Update generated files	a month ago
 coordination/v1beta1	Autogenerated stuff	2 months ago
 core/v1	Adding generated files	16 days ago
 events	Update generated files	a month ago
 extensions	Update generated files	27 days ago
 imagepolicy	regenerated all files and remove all YEAR fields	6 months ago
 networking	Update generated files	27 days ago
 policy	Update generated files	27 days ago
 rbac	Update generated files	27 days ago
 scheduling	Generated	3 months ago
 settings/v1alpha1	regenerated all files and remove all YEAR fields	6 months ago
 storage	Update generated files	27 days ago

k8s.io/api

- core/v1 →
- Pods
 - Services
 - ReplicaSet
 - ...

[Code](#)[Issues 8](#)[Pull requests 0](#)[Projects 0](#)[Wiki](#)[Insights](#)

No description, website, or topics provided.

[890 commits](#)[7 branches](#)[122 releases](#)[123 contributors](#)[Apache-2.0](#)

Branch: master ▾

[New pull request](#)[Create new file](#)[Upload files](#)[Find file](#)[Clone or download](#) ▾

k8s-publishing-bot Merge pull request #66252 from apelisse/dry-run ...	Latest commit cbafcd24 2 days ago
.github Treat staging repos as authoritative for all files	7 months ago
Godeps Merge pull request #65737 from roycaihw/api-linter	13 days ago
pkg dry-run: Run generated commands	10 days ago
third_party/forked/golang Fix error message in Equalities.DeepEqual	2 months ago
vendor Merge pull request #65034 from caesarxuchao/json-case-sensitive	a month ago
CONTRIBUTING.md Fix typo	6 months ago
LICENSE Add README and LICENSE to staging repos	9 months ago
OWNERS Name change: s/timstclair/tallclair/	a year ago
README.md apimachinery: fix typos in README	7 months ago
SECURITY_CONTACTS add PST to main SECURITY_CONTACTS as formality	2 months ago
code-of-conduct.md Add code-of-conduct.md to staging repos	7 months ago

- pkg/apis/meta/v1 →
- ObjectMeta
 - TypeMeta
 - ListOptions
 - DeleteOptions
 - GetOptions
 - Status
 - Event
 - ...

k8s.io/apimachinery

<https://github.com/openshift-talks/k8s-go/blob/master/client-go-basic/main.go>

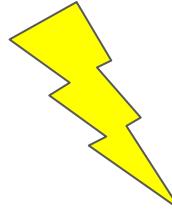
```
import (
    metav1 "k8s.io/apimachinery/pkg/apis/meta/v1"
    "k8s.io/client-go/tools/clientcmd"
    "k8s.io/client-go/kubernetes"
)

kubeconfig = flag.String("kubeconfig", "~/.kube/config", "path to the kubeconfig file")
flag.Parse()
config, err := clientcmd.BuildConfigFromFlags("", *kubeconfig)
clientset, err := kubernetes.NewForConfig(config)

pod, err := clientset.CoreV1().Pods("gophercon").Get("workshop", metav1.GetOptions{})
```

↖ *k8s.io/api/core/v1.Pod

↑
Get, Create, Update, List, Delete, Patch, Watch

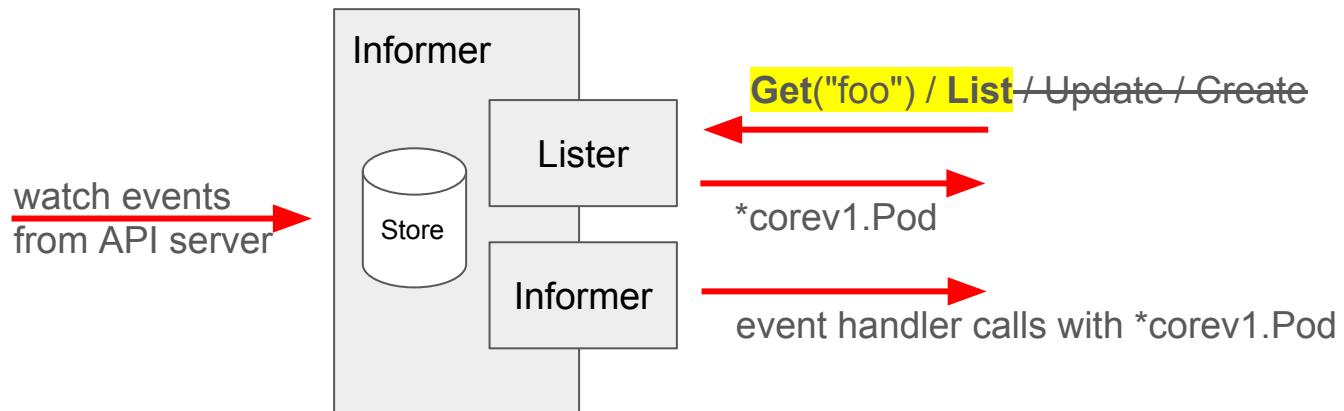


Katacoda

→ Client-go basics (step 3)

Informers

k8s.io/client-go/informers



Informers

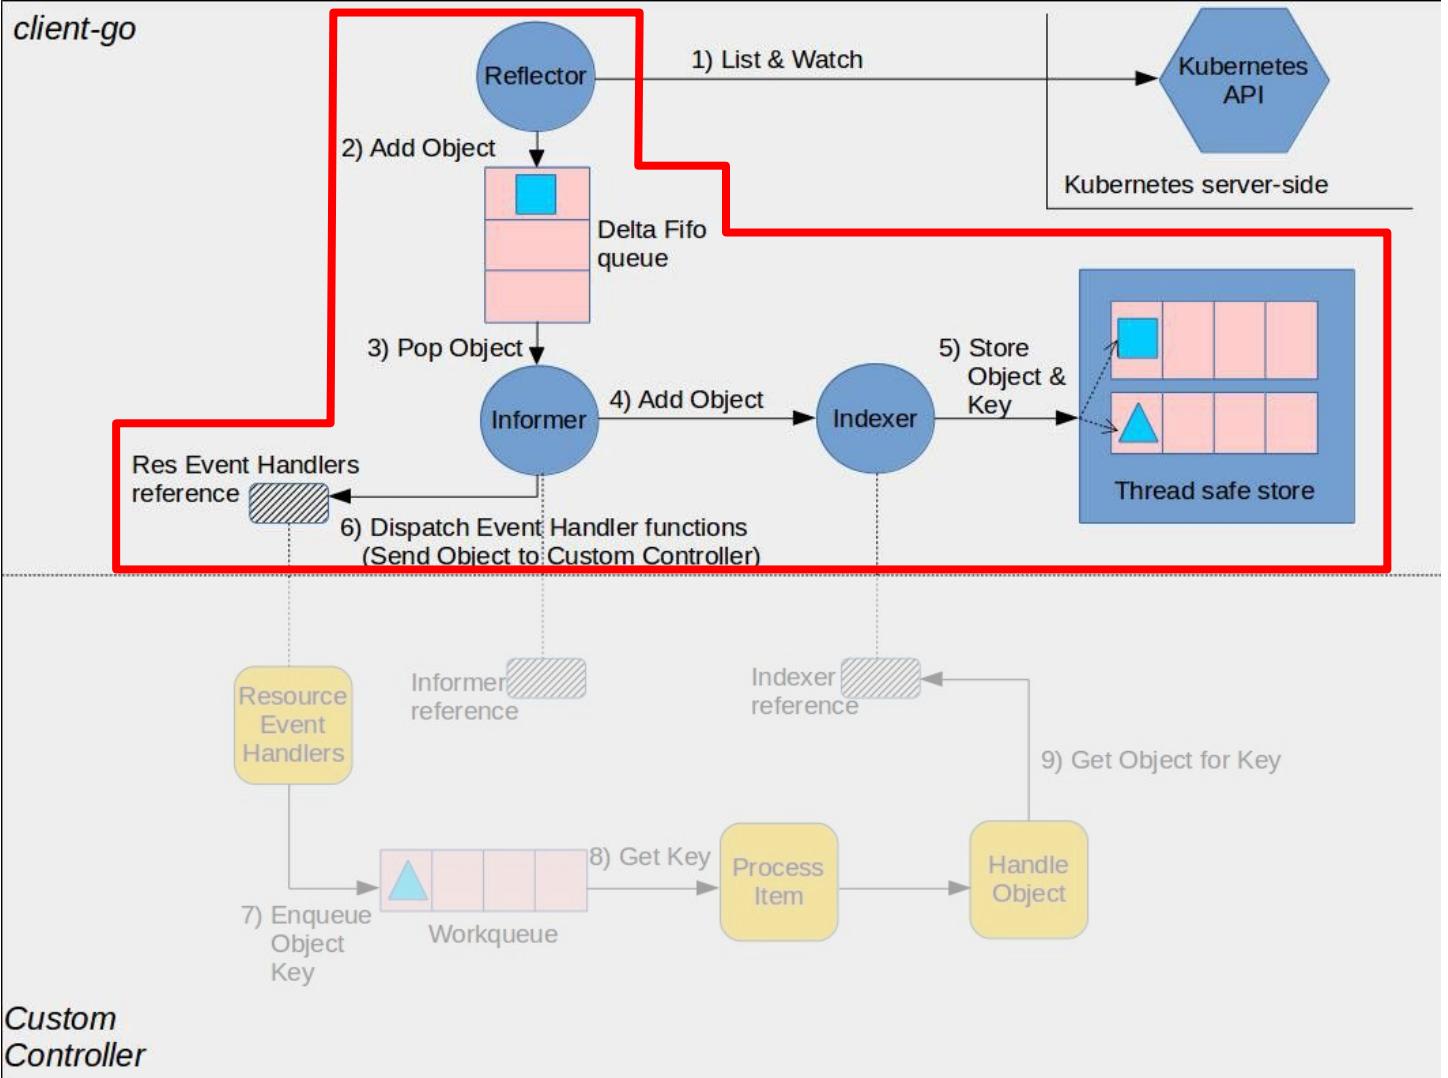


Image by @cloudark / Devdatta Kulkarni

<https://medium.com/@cloudark/kubernetes-custom-controllers-b6c7d0668fdf>

```
import (
    ...
    "k8s.io/client-go/informers"
)

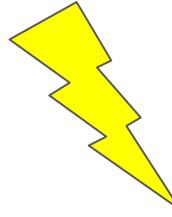
...
clientset, err := kubernetes.NewForConfig(config)
informerFactory := informers.NewSharedInformerFactory(clientset, time.Second*30)
podInformer := informerFactory.Core().V1().Pods()
podInformer.Informer().AddEventHandler(cache.ResourceEventHandlerFuncs{
    AddFunc: func(interface{}) {...},
    UpdateFunc: func(interface{}, interface{}) {...},
    DeleteFunc: func(interface{}) {...},
})
go informerFactory.Start(wait.NeverStop)

pod, err := podInformer.Lister().Pods("gophercon").Get("workshop")
```

*`k8s.io/api/core/v1.Pod`

in-memory cache – updated by List+Watch of the client

resync interval
= how often you re-see
your objects, even w/o
changes on the server



Katacoda

→ Client-go basics (step 3)

Vendoring client-go – versions and tags

k8s.io/client-go:

v7.0.0

k8s.io/api:

kubernetes-1.10.0

k8s.io/apimachinery:

kubernetes-1.10.0

v8.0.0

kubernetes-1.11.0

kubernetes-1.11.0

semver

no semver

Compatibility:

	Kubernetes 1.5	Kubernetes 1.6	Kubernetes 1.7	Kubernetes 1.8	Kubernetes 1.9	Kubernetes 1.10	Kubernetes 1.11
client-go 1.5	-	-	-	-	-	-	-
client-go 2.0	✓	+-	+-	+-	+-	+-	+-
client-go 3.0	+-	✓	-	+-	+-	+-	+-
client-go 4.0	+-	+-	✓	+-	+-	+-	+-
client-go 5.0	+-	+-	+-	✓	+-	+-	+-
client-go 6.0	+-	+-	+-	+-	✓	+-	+-
client-go 7.0	+-	+-	+-	+-	+-	✓	+-
client-go 8.0	+-	+-	+-	+-	+-	+-	✓
client-go HEAD	+-	+-	+-	+-	+-	+-	+-

Vendoring client-go – compatibility

- look at the **API groups** and **versions**: are group versions supported by the target cluster?
 - v1alpha1 – unstable:
 - might go away or change any time
 - often disabled by default
 - v1beta1 – towards stable:
 - exists at least one release in parallel to v1
 - no incompatible API change
 - v1 – stable, GA:
 - will stay
 - will be compatible
- look at the **Golang types** and their **fields**:
 - ⚠ we have **alpha fields** in v1, if alpha feature disabled:
some rejected, some ignored, some dropped

Vendoring client-go

- **source of truth:** Godeps/Godeps.json
- glide reads it
- dep only reads it on init, **not later on ensure!** 

=> specify all dependencies, at least k8s.io/{client-go,api,apimachinery}

dep Gopkg.toml

```
[[constraint]]
name = "k8s.io/api"
version = "kubernetes-1.10.0"

[[constraint]]
name = "k8s.io/apimachinery"
version = "kubernetes-1.10.0"

[[constraint]]
name = "k8s.io/client-go"
version = "7.0.0"

[prune]
go-tests = true
unused-packages = true
```



```
# the following overrides are necessary to enforce
# the given version, even though our
# code does not import the packages directly.

[[override]]
name = "k8s.io/api"
version = "kubernetes-1.10.0"

[[override]]
name = "k8s.io/apimachinery"
version = "kubernetes-1.10.0"

[[override]]
name = "k8s.io/client-go"
version = "7.0.0"
```

Kubernetes objects in Go

Golang Types

k8s.io/api/core/v1/types.go

```
type Pod struct {
    metav1.TypeMeta `json:",inline"`
    metav1.ObjectMeta `json:"metadata,omitempty"`

    Spec SessionSpec `json:"spec"`
    Status SessionStatus `json:"status"`
}
```

```
type PodSpec struct {
    Volumes []Volume `json:"volumes,omitempty"`
    Containers []Container `json:"containers"`
    NodeName string `json:"nodeName,omitempty"`
}
```

```
type PodStatus struct {
    Phase PodPhase `json:"phase,omitempty"`
    Conditions []PodCondition `json:"conditions,omitempty"`
    ContainerStatuses []ContainerStatus `json:"containerStatuses,omitempty"`
}
```

} user's desire

} how the controller saw it last time

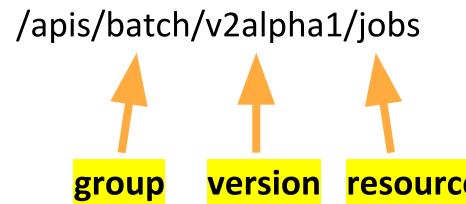
```
type TypeMeta struct {
    → Kind string `json:"kind,omitempty" protobuf:"bytes,1,opt,name=kind"`
    → APIVersion string `json:"apiVersion,omitempty" protobuf:"bytes,2,opt,name=apiVersion"`
}

type ObjectMeta struct {
    → Name string `json:"name,omitempty" protobuf:"bytes,1,opt,name=name"`
    → GenerateName string `json:"generateName,omitempty" protobuf:"bytes,2,opt,name=generateName"`
    → Namespace string `json:"namespace,omitempty" protobuf:"bytes,3,opt,name=namespace"`
    → SelfLink string `json:"selfLink,omitempty" protobuf:"bytes,4,opt,name=selfLink"`
    → UID types.UID `json:"uid,omitempty" protobuf:"bytes,5,opt,name=uid"`
    → ResourceVersion string `json:"resourceVersion,omitempty" protobuf:"bytes,6,opt,name=resourceVersion"`
    Generation int64 `json:"generation,omitempty" protobuf:"varint,7,opt,name=generation"`
    CreationTimestamp Time `json:"creationTimestamp,omitempty" protobuf:"timestamptime,8,opt,name=creationTimestamp"`
    DeletionTimestamp *Time `json:"deletionTimestamp,omitempty" protobuf:"timestamptime,9,opt,name=deletionTimestamp"`
    DeletionGracePeriodSeconds *int64 `json:"deletionGracePeriodSeconds,omitempty" protobuf:"varint,10,opt,name=deletionGracePeriodSeconds"`
    Labels map[string]string `json:"labels,omitempty" protobuf:"bytes,11,rep,name=labels"`
    Annotations map[string]string `json:"annotations,omitempty" protobuf:"bytes,12,rep,name=annotations"`
    OwnerReferences []OwnerReference `json:"ownerReferences,omitempty" protobuf:"bytes,13,rep,name=ownerReferences" patchStrategy:"merge"`
    Initializers *Initializers `json:"initializers,omitempty" protobuf:"bytes,14,opt,name=initializers"`
    Finalizers []string `json:"finalizers,omitempty" protobuf:"bytes,15,rep,name=finalizers" patchStrategy:"merge"`
    ClusterName string `json:"clusterName,omitempty" protobuf:"bytes,16,opt,name=clusterName"`
}
```

GVK and GVR

```
runtime.GroupVersionKind{Group: "batch", Version: "v2alpha1", Kind: "Job"}
```

```
runtime.GroupVersionResource{Group: "batch", Version: "v2alpha1", Resource: "jobs"}
```



Sometimes also: GroupVersion, GroupKind, GroupResource, ...

Golang types boilerplate

k8s.io/api/core/v1/register.go

```
const GroupName = "" // the legacy core group. Normally this would be e.g. apps.k8s.io
var SchemeGroupVersion = schema.GroupVersion{Group: GroupName, Version: "v1"}

func Kind(kind string) schema.GroupKind {
    return SchemeGroupVersion.WithKind(kind).GroupKind()
}

func Resource(resource string) schema.GroupResource {
    return SchemeGroupVersion.WithResource(resource).GroupResource()
}

var SchemeBuilder = runtime.NewSchemeBuilder(addKnownTypes)
var AddToScheme = SchemeBuilder.AddToScheme

// Adds the list of known types to Scheme.
func addKnownTypes(scheme *runtime.Scheme) error {
    scheme.AddKnownTypes(SchemeGroupVersion,
        &Pod{},
        &PodList{},
    )
    metav1.AddToGroupVersion(scheme, SchemeGroupVersion)
    return nil
}
```

← Install some generic types like Status,
GetOptions/ListOptions/DeleteOptions

Scheme – k8s.io/apimachinery/pkg/runtime.Scheme

k8s.io/client-go/kubernetes/scheme for a scheme with all Kube types

```
type Scheme struct {
    // versionMap allows one to figure out the go type of an object with
    // the given version and name.
    gvkToType map[schema.GroupVersionKind]reflect.Type

    // typeToGroupVersion allows one to find metadata for a given go object.
    // The reflect.Type we index by should *not* be a pointer.
    typeToGVK map[reflect.Type][]schema.GroupVersionKind

    // defaulterFuncs is an array of interfaces to be called with an object to provide defaulting
    // the provided object must be a pointer.
    defaulterFuncs map[reflect.Type]func(interface{})

    // converter stores all registered conversion functions. It also has
    // default converting behavior.
    converter *conversion.Converter
}

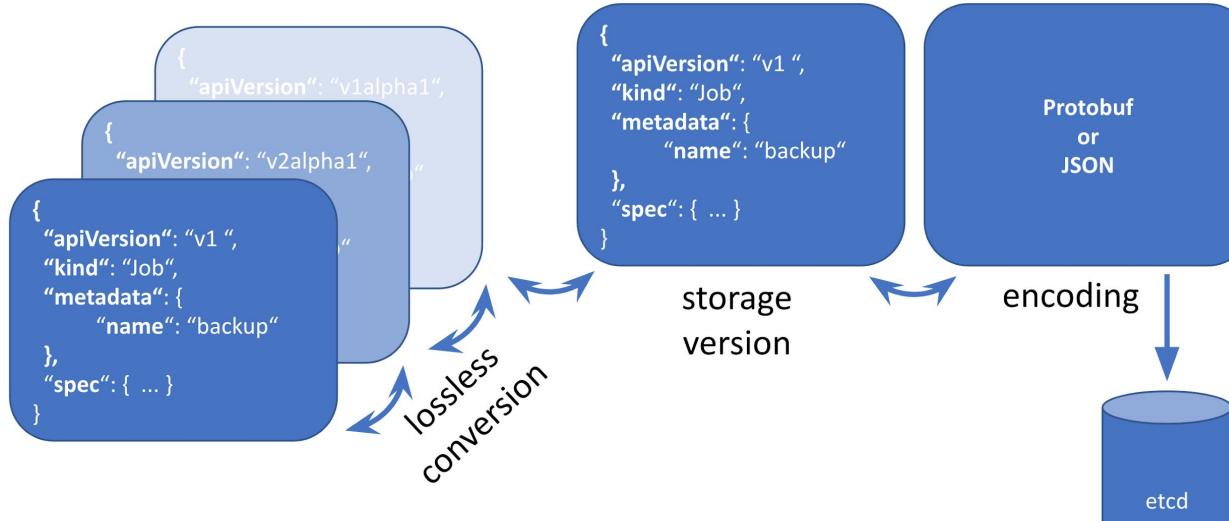
func (s *Scheme) AddKnownTypes(gv schema.GroupVersion, types ...Object) {
```

Conversion – scheme.Convert

```
func (s *Scheme) Convert(in, out interface{}, context interface{}) error
```

Only relevant in the API server implementation

- used by API servers internally to convert from e.g. **v1** objects to **v1beta1**
- ... usually by converting from **v1** to "**_internal**" to **v1beta**



Defaulting – `scheme.Default`

```
func (s *Scheme) Default(src Object)
```

Only relevant in the API server implementation

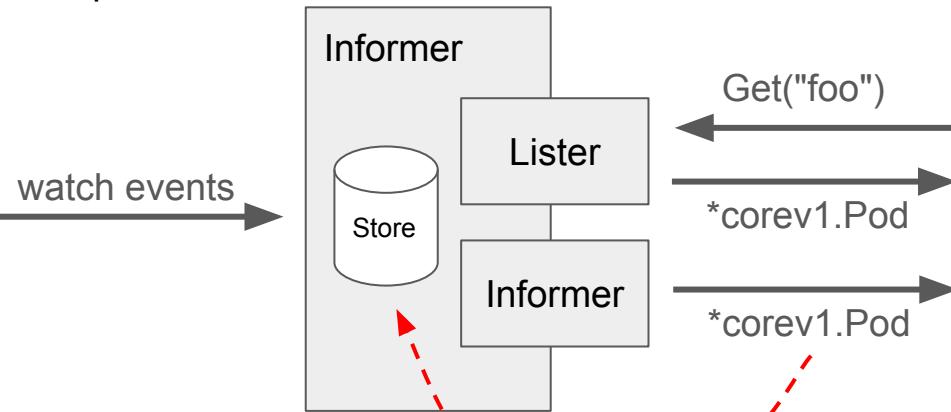
- done by the API server for all incoming objects and on read from etcd
 - sets fields that are unset**
- example: `PodSpec.Containers[*].RestartPolicy`

If unset, defaults to **RestartPolicyAlways**.

DeepCopy - obj.DeepCopy()

- generated for all types

Example:



Objects owned by the store! Never ever modify them!

```
func (in *Pod) DeepCopyInto(out *Pod) {  
    *out = *in  
    out.TypeMeta = in.TypeMeta  
    in.ObjectMeta.DeepCopyInto(&out.ObjectMeta)  
    in.Spec.DeepCopyInto(&out.Spec)  
    in.Status.DeepCopyInto(&out.Status)  
    return  
}  
  
func (in *Pod) DeepCopy() *Pod {  
    if in == nil {  
        return nil  
    }  
    out := new(Pod)  
    in.DeepCopyInto(out)  
    return out  
}  
  
func (in *Pod) DeepCopyObject() runtime.Object {  
    if c := in.DeepCopy(); c != nil {  
        return c  
    }  
    return nil  
}
```

RESTMapper

```
type RESTMapper interface {
    RESTMapping(gk schema.GroupKind, versions ...string) (*RESTMapping, error)
    KindFor(resource schema.GroupVersionResource) (schema.GroupVersionKind, error)
    ResourceFor(input schema.GroupVersionResource) (schema.GroupVersionResource, error)
}
```

\approx GroupVersionResource + cluster/namespace scope



Often used: DeferredDiscoveryRESTMapper with discovery

```
cachedClient := cacheddiscovery.NewMemCacheClient(clientset.Discovery())
restMapper := restmapper.NewDeferredDiscoveryRESTMapper(cachedClient)
```

Custom Resources (CRs) & Custom Resource Definitions (CRDs)

apiextensions/v1beta1

Custom Resource

example-healthcheckpolicy.yaml

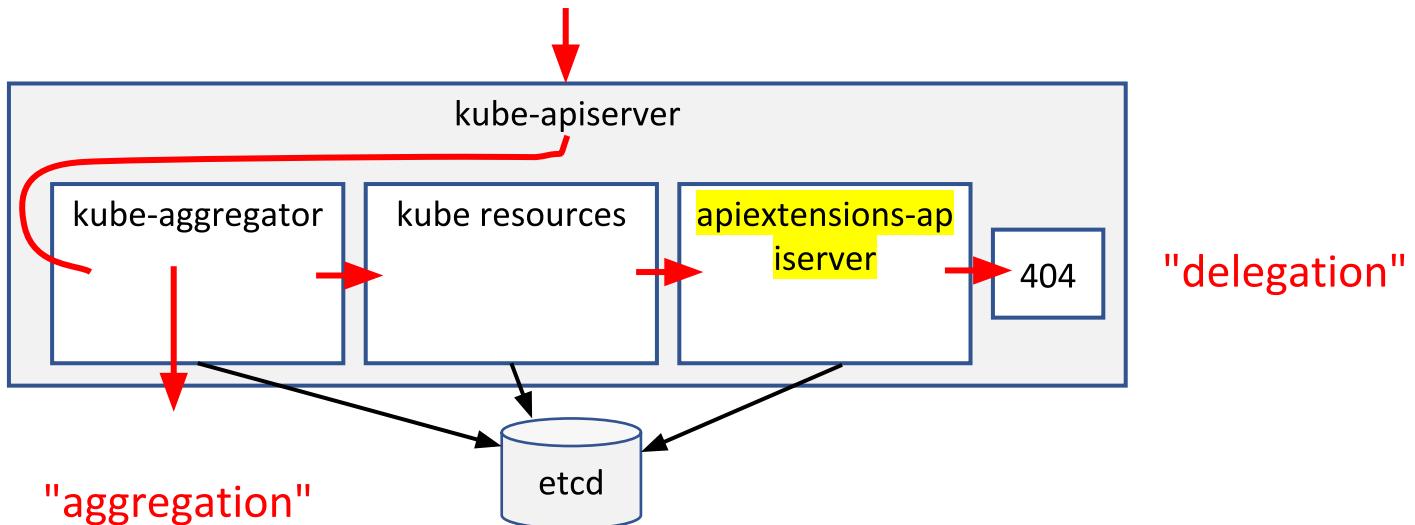
```
apiVersion: policy.k8s-go.openshift.org/v1alpha1
kind: HealthCheckPolicy
metadata:
  name: example
spec:
  podLabel: health-check-missing
status:
  statistics:
    pods:
      failed: 5
      passed: 42
  lastChecked: "So 29 Jul 2018 09:13:07 CEST"
```

We want to store these

apiextensions-apiserver inside kube-apiserver

<https://github.com/kubernetes/apiextensions-apiserver>

In every Kubernetes 1.7+ cluster.



Custom Resource

example-healthcheckpolicy.yaml

```
apiVersion: policy.k8s-go.openshift.org/v1alpha1
kind: HealthCheckPolicy
metadata:
  name: example
spec:
  podLabel: health-check-missing
status:
  statistics:
    pods:
      failed: 5
      passed: 42
  lastChecked: "So 29 Jul 2018 09:13:07 CEST"
```

We want to store these

Custom Resource Definition (CRD)

apiextensions/v1beta1

Defines how CRs are stored

```
apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  name: healthcheckpolicies.policy.k8s-go.openshift.org
spec:
  group: policy.k8s-go.openshift.org
  version: v1alpha1
  names:
    kind: HealthCheckPolicy
    plural: healthcheckpolicies
    scope: Namespaced
```

The diagram shows two orange arrows pointing upwards from the 'names' section of the CRD spec to the 'group' and 'plural' fields in the 'spec' section. A blue curved arrow points from the text 'must match' to the word 'names'.

must match

Custom Resource Definition (CRD)

apiextensions/v1beta1

```
$ kubectl create -f healthcheckpolicy-crd.yaml
```

```
apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  name: healthcheckpolicies.policy.k8s-go.openshift.org
spec:
  group: policy.k8s-go.openshift.org
  version: v1alpha1
  names:
    kind: HealthCheckPolicy
    plural: healthcheckpolicies
    scope: Namespaced
```

```
...  
status:  
acceptedNames:  
  kind: HealthCheckPolicy  
  listKind: HealthCheckPolicyList  
  plural: healthcheckpolicies  
  singular: healthcheckpolicy  
conditions:
```

```
- type: NamesAccepted ←  
  message: no conflicts found  
  reason: NoConflicts  
  status: "True"  
- type: Established ←  
  message: the initial names have  
  been accepted  
  reason: InitialNamesAccepted  
  status: "True"
```

```
$ kubectl get healthcheckpolicies -v=7
```

- I0429 21:17:53.042783 66743 round_trippers.go:383] GET <https://localhost:6443/apis>
- I0429 21:17:53.135811 66743 round_trippers.go:383] GET <https://localhost:6443/apis/policy.k8s-go.openshift.org/v1alpha1>
- I0429 21:17:53.138353 66743 round_trippers.go:383] GET <https://localhost:6443/apis/policy.k8s-go.openshift.org/v1alpha1/namespaces/default/healthcheckpolicies>

No resources found.

healthcheckpolicies → kind HealthCheckPolicy
resource healthcheckpolicies



We call this "REST mapping"

```
$ http localhost:8080/apis/
{
    "groups": [
        {
            "name": "policy.k8s-go.openshift.org",
            "preferredVersion": {"groupVersion": "policy.k8s-go.openshift.org/v1", "version": "v1alpha1"},
            "versions": [{"groupVersion": "policy.k8s-go.openshift.org/v1alpha1", "version": "v1alpha1"}]
        }, ...
    }
}
```

```
$ http localhost:8080/apis/policy.k8s-go.openshift.org/v1alpha1
```

```
{
    "apiVersion": "v1",
    "groupVersion": "policy.k8s-go.openshift.org/v1alpha1",
    "kind": "APIResourceList",
    "resources": [
        {
            "kind": "HealthCheckPolicy",
            "name": "healthcheckpolicies",
            "namespaced": true,
            "verbs": ["create", "delete", "deletecollection",
                      "get", "list", "patch", "update", "watch"
            ]
        }, ...
    }
}
```

resource name ⇒  /apis/policy.k8s-go.openshift.org/v1alpha1/healthcheckpolicies

```
$ kubectl get healthcheckpolicies -v=7
```

- I0429 21:17:53.042783 66743 round_trippers.go:383] GET <https://localhost:6443/apis>
- I0429 21:17:53.135811 66743 round_trippers.go:383] GET <https://localhost:6443/apis/policy.k8s-go.openshift.org/v1alpha1>
- I0429 21:17:53.138353 66743 round_trippers.go:383] GET <https://localhost:6443/apis/policy.k8s-go.openshift.org/v1alpha1/namespaces/default/healthcheckpolicies>

No resources found.

healthcheckpolicies → kind HealthCheckPolicy
resource healthcheckpolicies



We call this "REST mapping"

```
$ kubectl get healthcheckpolicies -v=7
```

- I0429 21:17:53.042783 66743 round_trippers.go:383] GET <https://localhost:6443/apis>
- I0429 21:17:53.135811 66743 round_trippers.go:383] GET <https://localhost:6443/apis/policy.k8s-go.openshift.org/v1alpha1>
- I0429 21:17:53.138353 66743 round_trippers.go:383] GET <https://localhost:6443/apis/policy.k8s-go.openshift.org/v1alpha1/namespaces/default/healthcheckpolicies>

No resources found.

hcp
note: a "shortName"

→ kind HealthCheckPolicy
resource healthcheckpolicies

We call this "REST mapping"



1

Create the
CRD

```
$ kubectl create -f healthcheckpolicies-crd.yaml
```

```
apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  name: healthcheckpolicies.policy.k8s-go...
spec:
  group: policy.k8s-go.openshift.org
  version: v1alpha1
  names:
    kind: HealthCheckPolicy
    plural: healthcheckpolicies
    scope: Namespaced
```

2

Create
CustomResources

```
$ kubectl create -f example-healthcheckpolicy.yaml
```

```
apiVersion:
  policy.k8s-go.openshift.org/v1alpha1
kind: HealthCheckPolicy
metadata:
  name: example
spec:
  podLabel: health-check-missing
status:
  statistics:
    pods:
      failed: 5
      passed: 42
  lastChecked: "So 29 Jul 2018 09:13:07 CEST"
```

1

Create the
CRD

```
$ kubectl create -f healthcheckpolicies-crd.yaml
```

```
apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  name: healthcheckpolicies.policy.k8s-go...
spec:
  group: policy.k8s-go.openshift.org
  version: v1alpha1
  names:
    kind: HealthCheckPolicy
    plural: healthcheckpolicies
    scope: Namespaced
```

2

Create
CustomResources

```
$ kubectl create -f example-healthcheckpolicy.yaml
```

```
apiVersion:
  policy.k8s-go.openshift.org/v1alpha1
kind: HealthCheckPolicy
metadata:
  name: example
spec:
  podLabel: health-check-missing
status:
  statistics:
    pods:
      failed: 5
      passed: 42
  lastChecked: "So 29 Jul 2018 09:13:07 CEST"
```

Follow spec+status pattern.

Watch, i.e. event stream

```
$ kubectl get healthcheckpolicies -w --no-headers
example <none> {"apiVersion":"policy.k8s-go.openshift.org/v1","kind":"HealthCheckPolicies",...
example2 <none> {"apiVersion":"policy.k8s-go.openshift.org/v1","kind": " HealthCheckPolicies ",...}
```

```
$ curl -f
'http://127.0.0.1:8080/apis/policy.k8s-go.openshift.org/v1/namespaces/default/healthcheckpolicies?watch=true&resourceVersion=434'
```

```
{"type":"DELETED","object":{"apiVersion":"policy.k8s-go.openshift.org/v1","kind":"HealthCheckPolicies","metadata":{"name":"example","namespace":"default","selfLink":"/apis/policy.k8s-go.openshift.org/v1/namespaces/default/healthcheckpolicies/example","uid":"8f5312c0-29c8-11e7-88f9-4c3275978b79","resourceVersion":"435","creationTimestamp":"2017-04-25T15:05:03Z"},"spec":{...}}}
```

```
{"type":"ADDED","object":{"apiVersion":"policy.k8s-go.openshift.org/v1","kind":"HealthCheckPolicies","metadata":{"name":"example2","namespace":"default","selfLink":"/apis/policy.k8s-go.openshift.org/v1/namespaces/default/healthcheckpolicies/example2","uid":"b4318cb5-29c8-11e7-88f9-4c3275978b79","resourceVersion":"436","creationTimestamp":"2017-04-25T15:06:05Z"},"spec":{...}}}
```

Validation

- The standard: [OpenAPI v3 schema](#)
- based on [JSON Schema](#)

```
apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  name: healthcheckpolicies.policy.k8s-go...
spec:
  group: policy.k8s-go.openshift.org
  version: v1alpha1
  names: ...
  validation:
    openAPIV3Schema: <see next slide>
```

Custom Resource

spec:

action: label

podLabel: health-check-missing

status:

statistics:

pods:

failed: 5

passed: 42

lastChecked: "So 29 Jul 2018 09:13:07..."

properties:

spec:

properties:

action:

anyOf: [{"pattern": "^label|kill|annotate\$"}, ...]

default: "label"

a quantor (anyOf, oneOf, allOf
exist)

note: enum is forbidden (why?)

podLabel: {"type": "string", "pattern": "^[0-9]+-...\$"}
required: ["action"]

probably in 1.12+

status:

properties:

statistics:

properties:

failed: {"type": "integer", "minimum": "0"}

passed: {"type": "integer", "minimum": "0"}

required: []

regular
expression

Helpful tools:

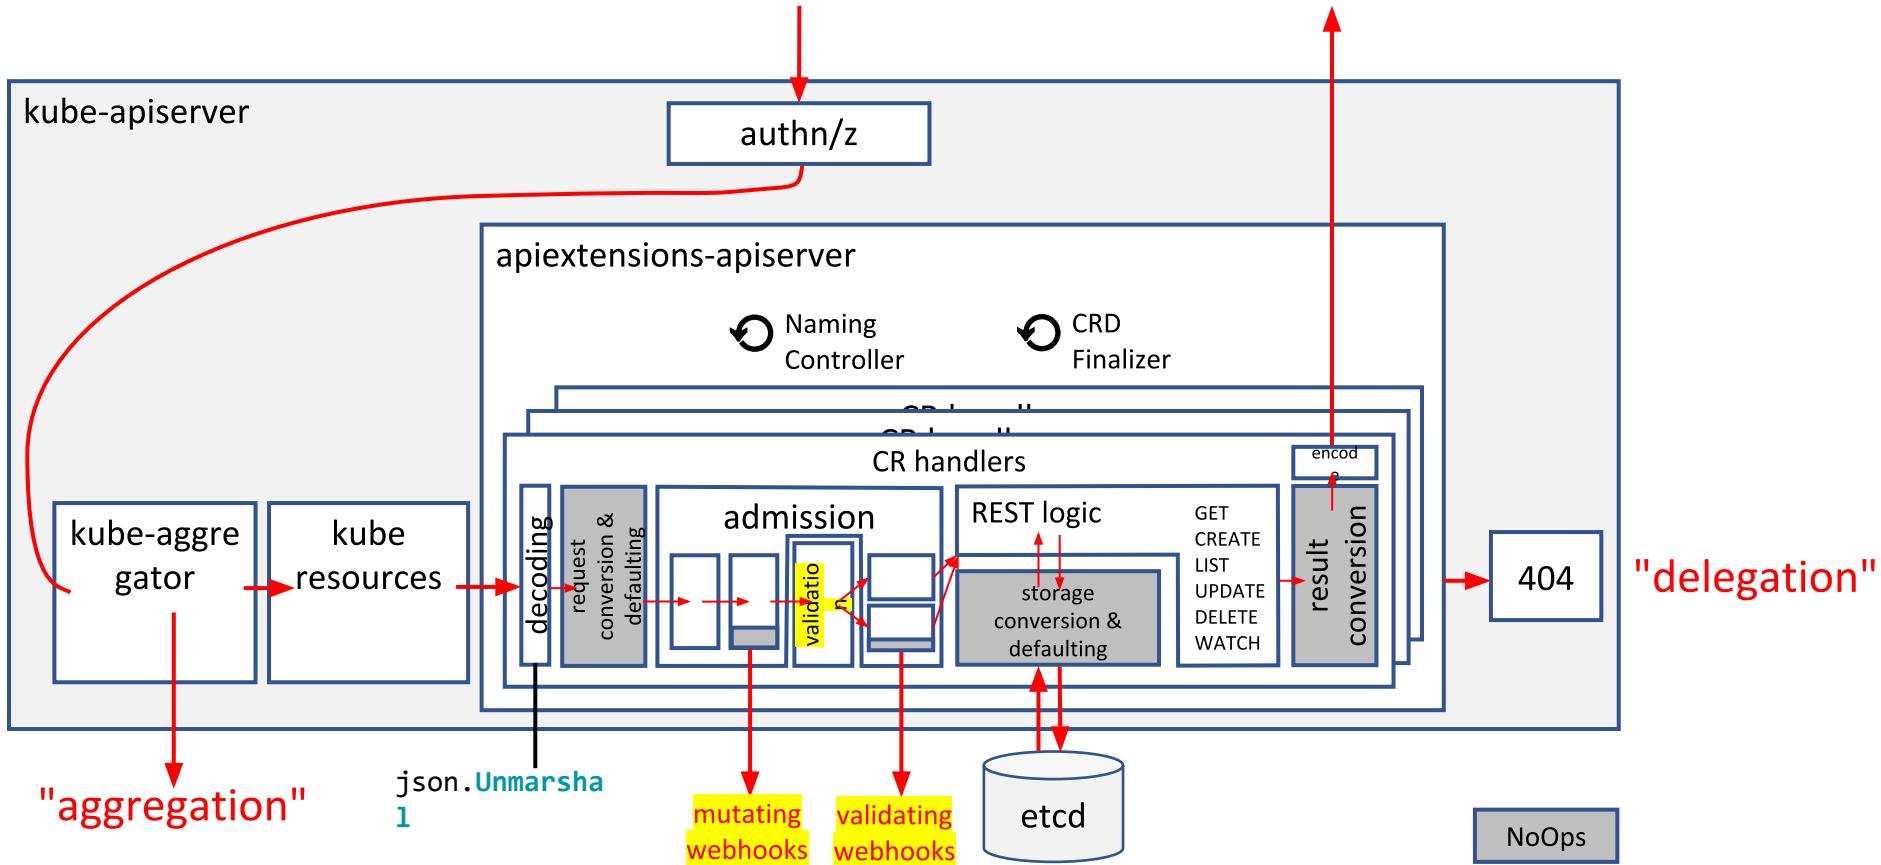
[kubernetes/kube-openapi#37](#)

[tamalsaha/kube-openapi-generator](#)

Some other tool from prometheus-operator?

Rancher has another one, speak to @lemonjet

Zoom into apiextensions-apiserver



"CRDs are limited"

- ~~no version conversion (only one version possible per CRD)~~
 - in 1.11 **multiple versions**, but no conversion
 - in 1.12+ **conversions** (via webhook) come, but slowly (i.e. first alpha)
- no **defaulting** (maybe in 1.12 as alpha)
- no **validation** (beta in 1.9, Google Summer of Code project)
- ~~no subresources~~ (scale, status) (beta in 1.11)
- ~~no admission~~ (since 1.7: admission webhooks + initializers)
- alpha CRDs are beta
- no **custom printing** in kubectl (in 1.11: custom printer columns)



Katacoda

→ CRDs for users (step 4)

Towards a controller: Golang types

pkg/apis/policy/v1alpha1/types.go

```
type HealthCheckPolicy struct {
    metav1.TypeMeta `json:",inline"`
    metav1.ObjectMeta `json:"metadata,omitempty"`

    Spec HealthCheckPolicySpec `json:"spec"`
    Status HealthCheckPolicyStatus `json:"status"`
}

type HealthCheckPolicySpec struct {
}

type HealthCheckPolicyStatus struct {
    PodsFailed int64 `json:"podsFailed,omitempty"`
}
```

Towards a controller: Golang types

pkg/apis/policy/v1alpha1/types.go

```
type HealthCheckPolicy struct {
    metav1.TypeMeta `json:",inline"`
    metav1.ObjectMeta `json:"metadata,omitempty"`

    Spec HealthCheckPolicySpec `json:"spec"`
    Status HealthCheckPolicyStatus `json:"status"`
}

type HealthCheckPolicySpec struct {
}

type HealthCheckPolicyStatus struct {
    PodsFailed int64 `json:"podsFailed,omitempty"`
}
```

doc.go:

```
// +k8s:deepcopy-gen=package
package v1
```

More info about code generation in
my OpenShift blog post
[Code Generation for
CustomResources](#)

Towards a controller: Golang types

pkg/apis/policy/v1alpha1/types.go

```
// +k8s:deepcopy-gen:interfaces=k8s.io/apimachinery/pkg/runtime.Object
type HealthCheckPolicy struct {
    metav1.TypeMeta `json:",inline"`
    metav1.ObjectMeta `json:"metadata,omitempty"`

    Spec  HealthCheckPolicySpec `json:"spec"`
    Status HealthCheckPolicyStatus `json:"status"`
}

type HealthCheckPolicySpec struct {

}

type HealthCheckPolicyStatus struct {
    PodsFailed int64 `json:"podsFailed,omitempty"`
}
```

doc.go:

```
// +k8s:deepcopy-gen=package
package v1
```

More info about code generation in
my OpenShift blog post
[Code Generation for
CustomResources](#)

Interface runtime.Object

```
type Object interface {
    GetObjectKind() schema.ObjectKind
    DeepCopyObject() Object
}
```

= everything that has TypeMeta

```
func (p *Pod) DeepCopyObject() runtime.Object {
    if c := in.DeepCopy(); c != nil {
        return c
    }
    return nil
}
```

// Go riddle: why?

generated through: // +k8s:deepcopy-gen:interfaces=k8s.io/apimachinery/pkg/runtime.Object

Towards a controller: Golang types

pkg/apis/policy/v1alpha1/types.go

```
// +k8s:deepcopy-gen:interfaces=k8s.io/apimachinery/pkg/runtime.Object
type HealthCheckPolicy struct {
    metav1.TypeMeta `json:",inline"`
    metav1.ObjectMeta `json:"metadata,omitempty"`

    Spec  HealthCheckPolicySpec `json:"spec"`
    Status HealthCheckPolicyStatus `json:"status"`
}

type HealthCheckPolicySpec struct {

}

type HealthCheckPolicyStatus struct {
    PodsFailed int64 `json:"podsFailed,omitempty"`
}
```

doc.go:

```
// +k8s:deepcopy-gen=package
package v1
```

More info about code generation in
my OpenShift blog post
[Code Generation for
CustomResources](#)

Towards a controller: Golang types

pkg/apis/policy/v1alpha1/types.go

```
// +genclient
// +k8s:deepcopy-gen:interfaces=k8s.io/apimachinery/pkg/runtime.Object
type HealthCheckPolicy struct {
    metav1.TypeMeta `json:",inline"`
    metav1.ObjectMeta `json:"metadata,omitempty"`

    Spec  HealthCheckPolicySpec `json:"spec"`
    Status HealthCheckPolicyStatus `json:"status"`
}

type HealthCheckPolicySpec struct {

}

type HealthCheckPolicyStatus struct {
    PodsFailed int64 `json:"podsFailed,omitempty"`
}
```

doc.go:

```
// +k8s:deepcopy-gen=package
package v1
```

More info about code generation in
my OpenShift blog post
[Code Generation for
CustomResources](#)

Towards a controller: Golang types

pkg/apis/policy/v1alpha1/types.go

```
// +genclient
// +k8s:deepcopy-gen:interfaces=k8s.io/apimachinery/pkg/runtime.Object
type HealthCheckPolicy struct {
    metav1.TypeMeta `json:",inline"`
    metav1.ObjectMeta `json:"metadata,omitempty"`

    Spec HealthCheckPolicySpec `json:"spec"`
    Status HealthCheckPolicyStatus `json:"status"`
}
```

doc.go:

```
// +k8s:deepcopy-gen=package
package v1
```

```
// +k8s:deepcopy-gen:interfaces=k8s.io/apimachinery/pkg/runtime.Object
type HealthCheckPolicyList struct {
    metav1.TypeMeta `json:",inline"`
    metav1.ListMeta `json:"metadata"`

    Items []HealthCheckPolicy `json:"items"`
}

type HealthCheckPolicySpec struct {
```

More info about code generation in
my OpenShift blog post
[Code Generation for
CustomResources](#)

Towards a controller: Golang types

pkg/apis/policy/v1alpha1/register.go

```
const GroupName = "policy.k8s-go.openshift.org"

var SchemeGroupVersion = schema.GroupVersion{Group: GroupName, Version: "v1alpha1"}

func Kind(kind string) schema.GroupKind {
    return SchemeGroupVersion.WithKind(kind).GroupKind()
}

func Resource(resource string) schema.GroupResource {
    return SchemeGroupVersion.WithResource(resource).GroupResource()
}

var SchemeBuilder = runtime.NewSchemeBuilder(addKnownTypes)
var AddToScheme = SchemeBuilder.AddToScheme

// Adds the list of known types to Scheme.
func addKnownTypes(scheme *runtime.Scheme) error {
    scheme.AddKnownTypes(SchemeGroupVersion,
        &HealthCheckPolicy{},
        &HealthCheckPolicies{},
    )
    metav1.AddToGroupVersion(scheme, SchemeGroupVersion)
    return nil
}
```

hint: copy from some other API group

Code Generation

all =
deepcopy,defaulter,client,lister,informer

```
$ vendor/k8s.io/code-generator/generate-groups.sh all \  
github.com/openshift/k8s-go/crd-go/pkg/generated \  
github.com/openshift/k8s-go/crd-go/pkg/pkg/apis \  
"policy:v1alpha1,v2,v3 someothergroup.io:v1" \  
--output-base "$GOPATH/src" \  
--go-header-file hack/boilerplate.go.txt
```

Usually put into [hack/update-codegen.sh](#) (compare [k8s.io/sample-controller](#)).

doc.go for crazy group and package names:

```
// +groupName=example.test.crd.code-generator.k8s.io  
// +groupGoName=SecondExample
```

Alternatives later: kubebuilder, operator-sdk

Native CRD Client

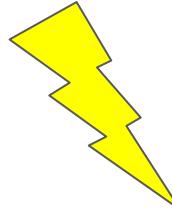
```
import (
    metav1 "k8s.io/apimachinery/pkg/apis/meta/v1"
    "k8s.io/client-go/tools/clientcmd"
    policyclient "github.com/openshift/k8s-go/crd-go/pkg/generated/clientset/versioned"
)

kubeconfig = flag.String("kubeconfig", "~/.kube/config", "path to the kubeconfig file")
flag.Parse()
config, err := clientcmd.BuildConfigFromFlags("", *kubeconfig)
clientset, err := kubecon.NewForConfig(config)

policy, err := clientset.PolicyV1alpha1().HealthCheckPolicies("default").Get("example", {})
```

similarly `/informers` and `/listers`

`*github.com/openshift/k8s-go/crd-go/pkg/apis/policy/v1alpha1.HealthCheckPolicy`



Katacoda

→ CRDs for developers (step 5)

Outlook – Custom Resources

- Kubernetes 1.11+
 - α: **Multiple versions without conversion** – [design proposal](#)
 - α: **Server Side Printing Columns** – “kubectl get” customization – [#60991](#)
 - β: **Subresources** – α since 1.10 – [#62786](#)
 - OpenAPI **additionalProperties** allowed now
(mutually exclusive with properties)
- Kubernetes 1.12
 - α: **Conversion Webhooks**
 - α: **Pruning** – in validation spec unspecified fields are removed – **blocker for GA**
 - α: **Defaulting** – defaults from OpenAPI validation schema are applied

Controllers

Operator

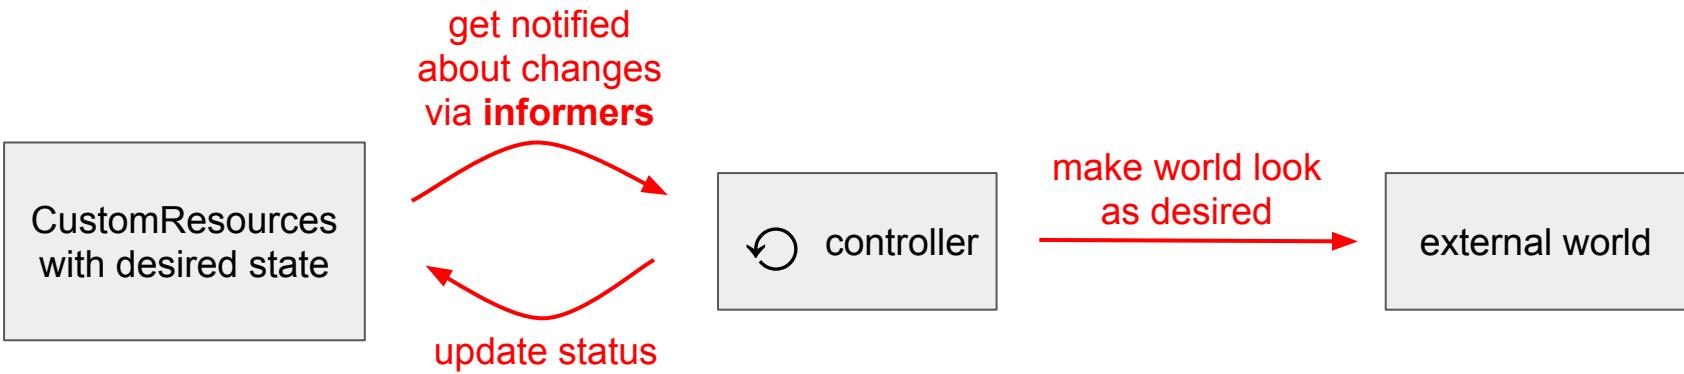
Controller

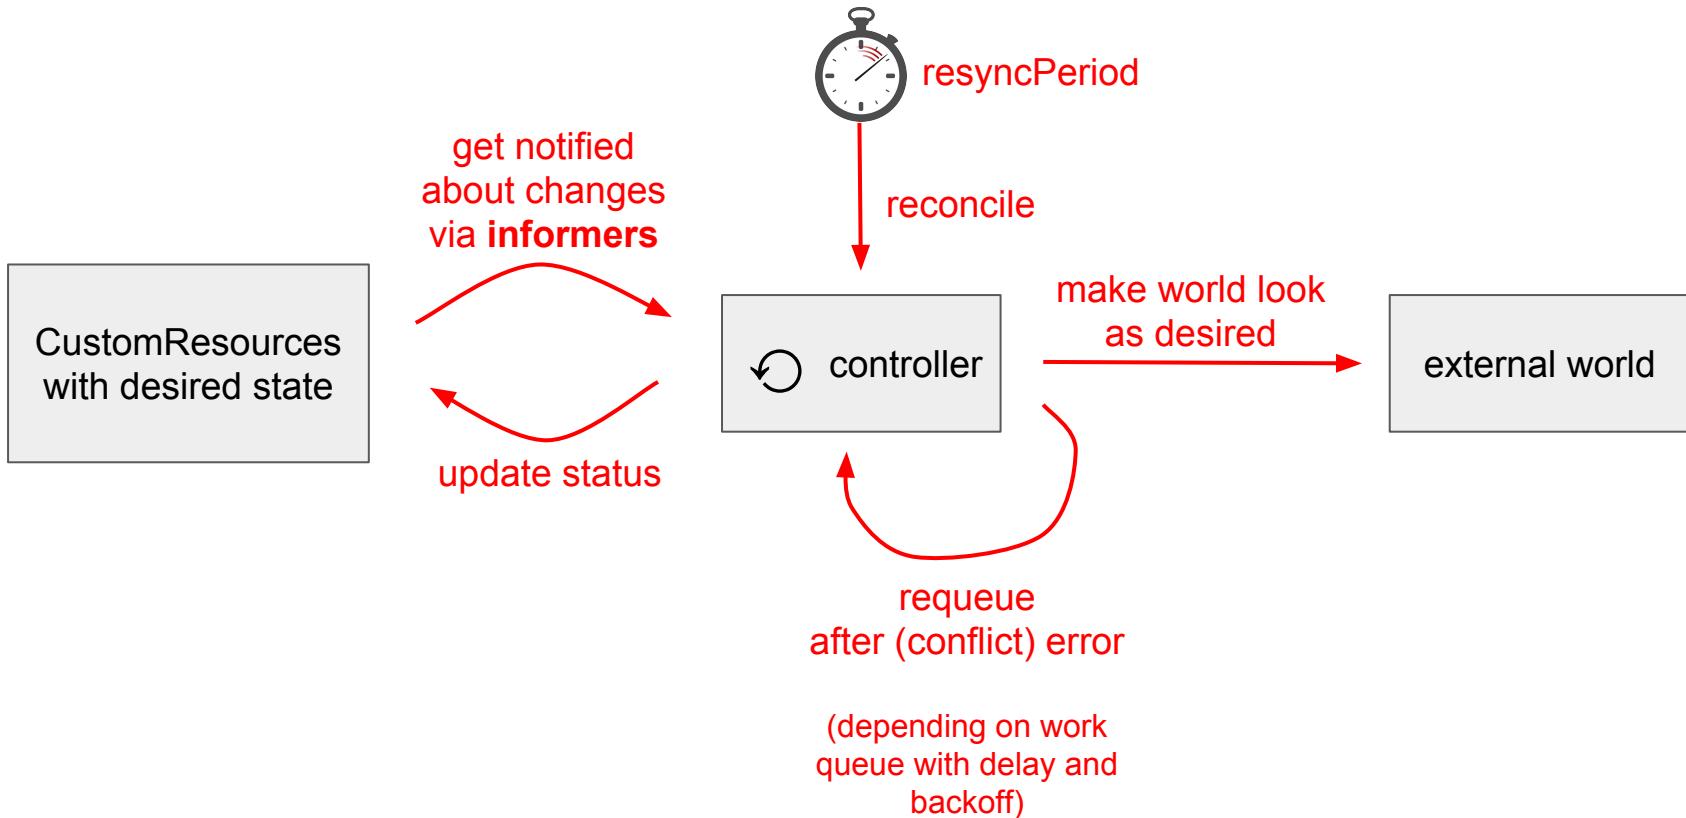
operates service with
special operational
knowledge?

just updates some
Kubernetes object

1. **read** object (preferably event driven with **watches**)
2. **change** object, update the world
3. **update** object on apiserver
4. **repeat**







Inner workings of a Controller

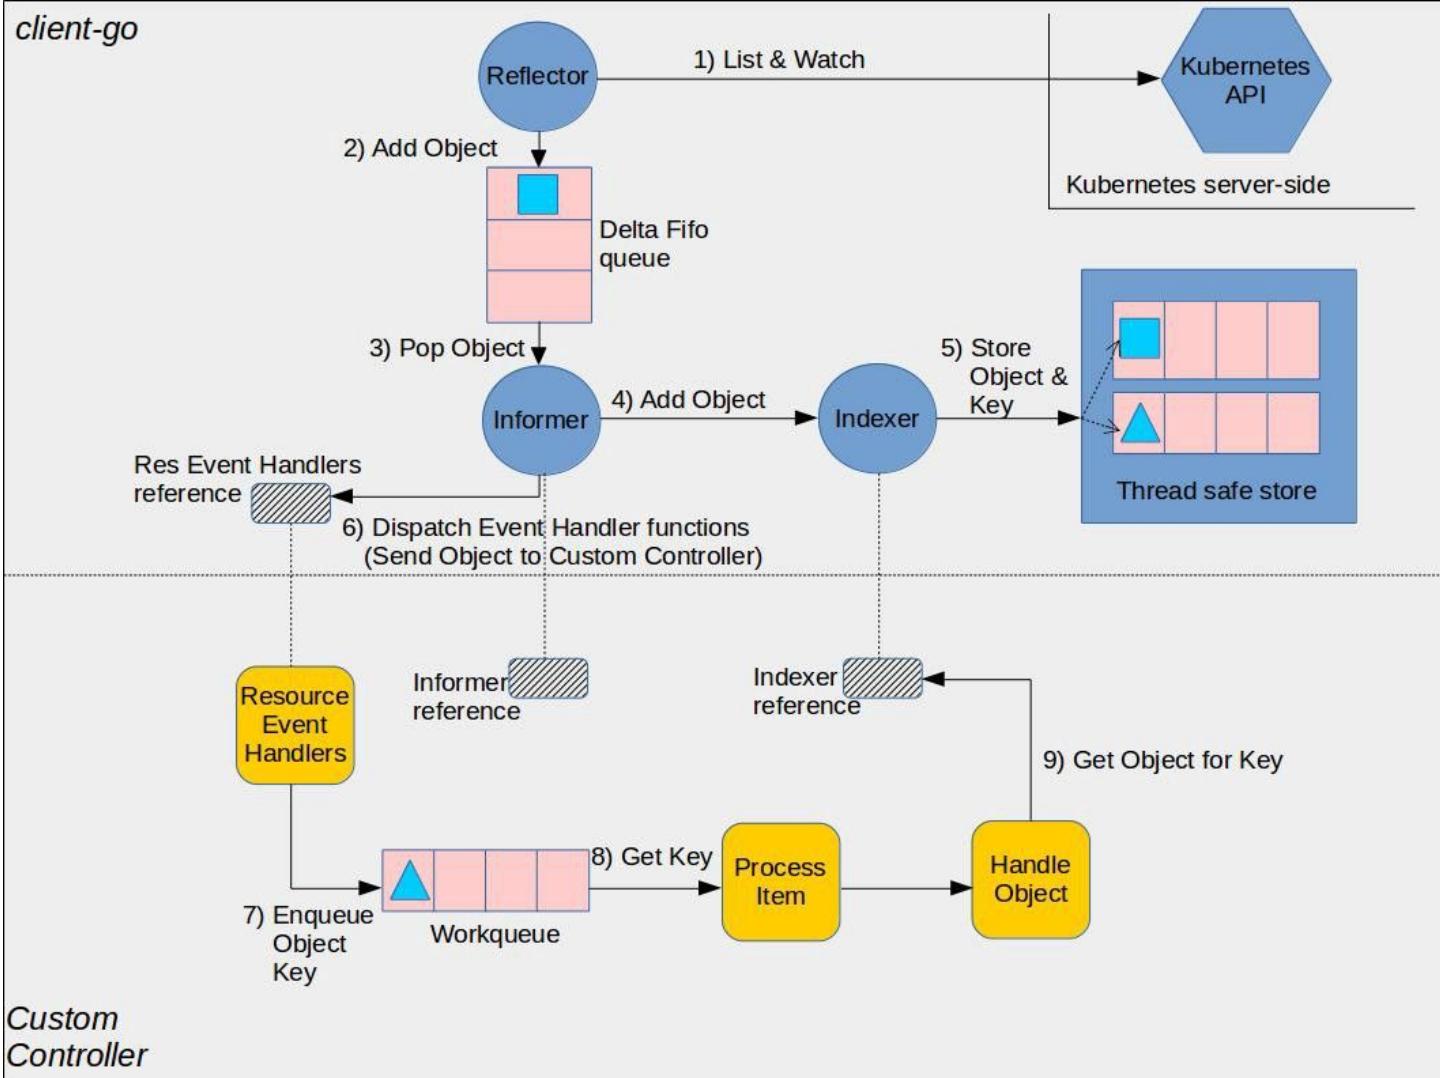


Image by @cloudark / Devdatta Kulkarni

<https://medium.com/@cloudark/kubernetes-custom-controllers-b6c7d0668fdf>

Inner workings of a Controller

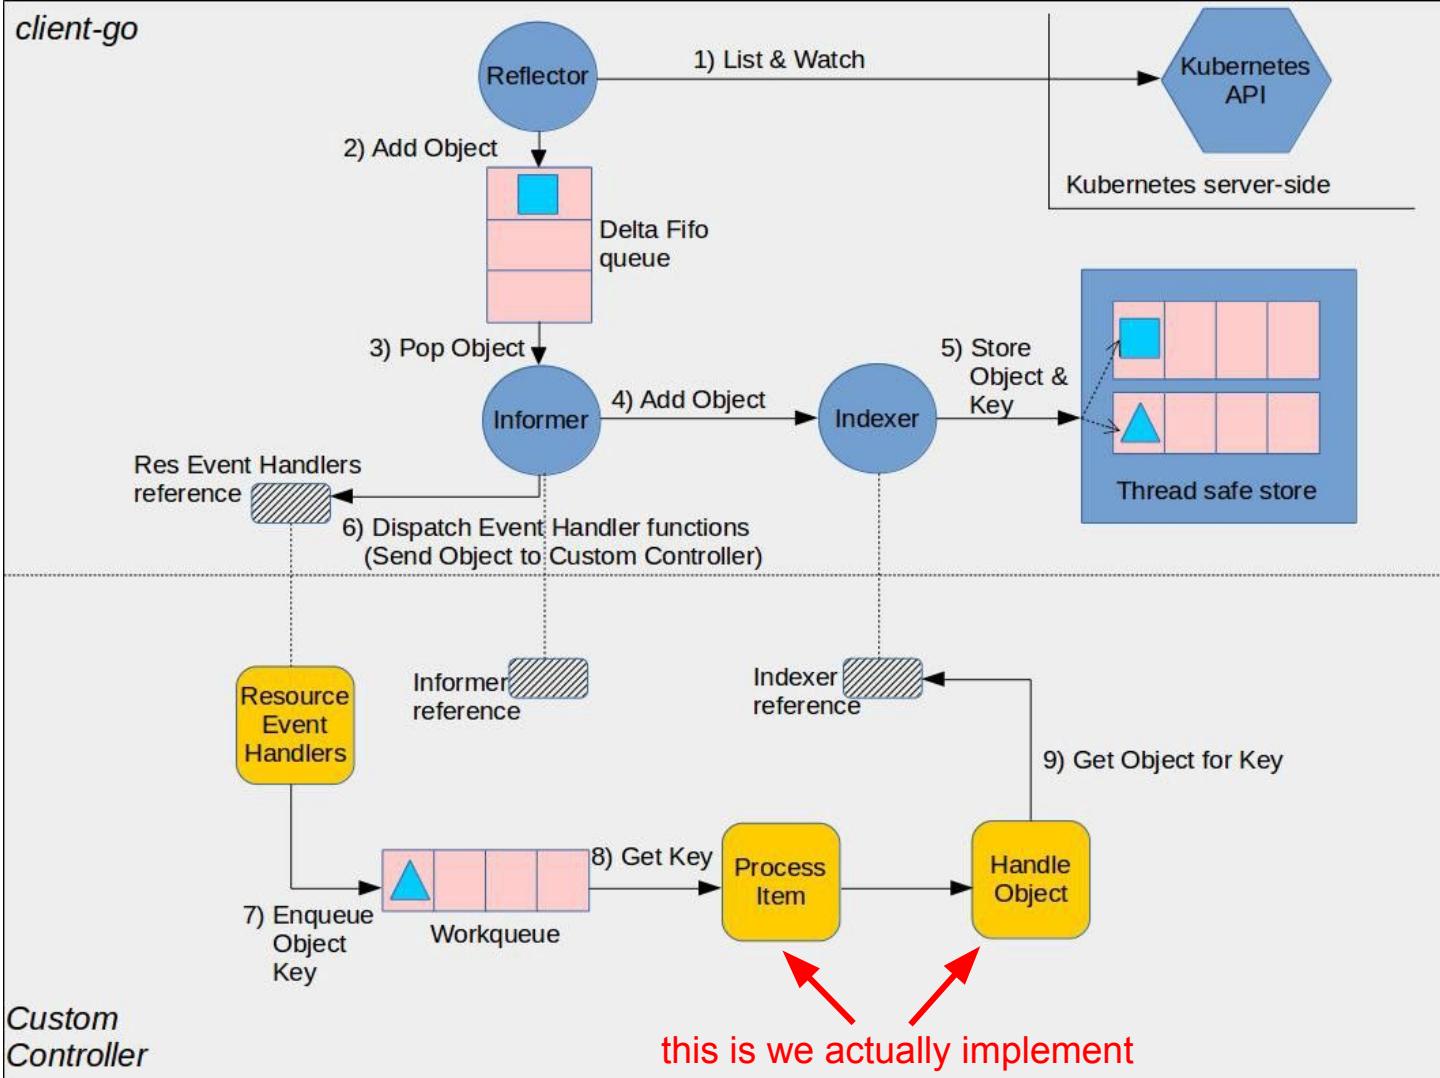


Image by @cloudark / Devdatta Kulkarni

<https://medium.com/@cloudark/kubernetes-custom-controllers-b6c7d0668fdf>

Inner workings of a Controller

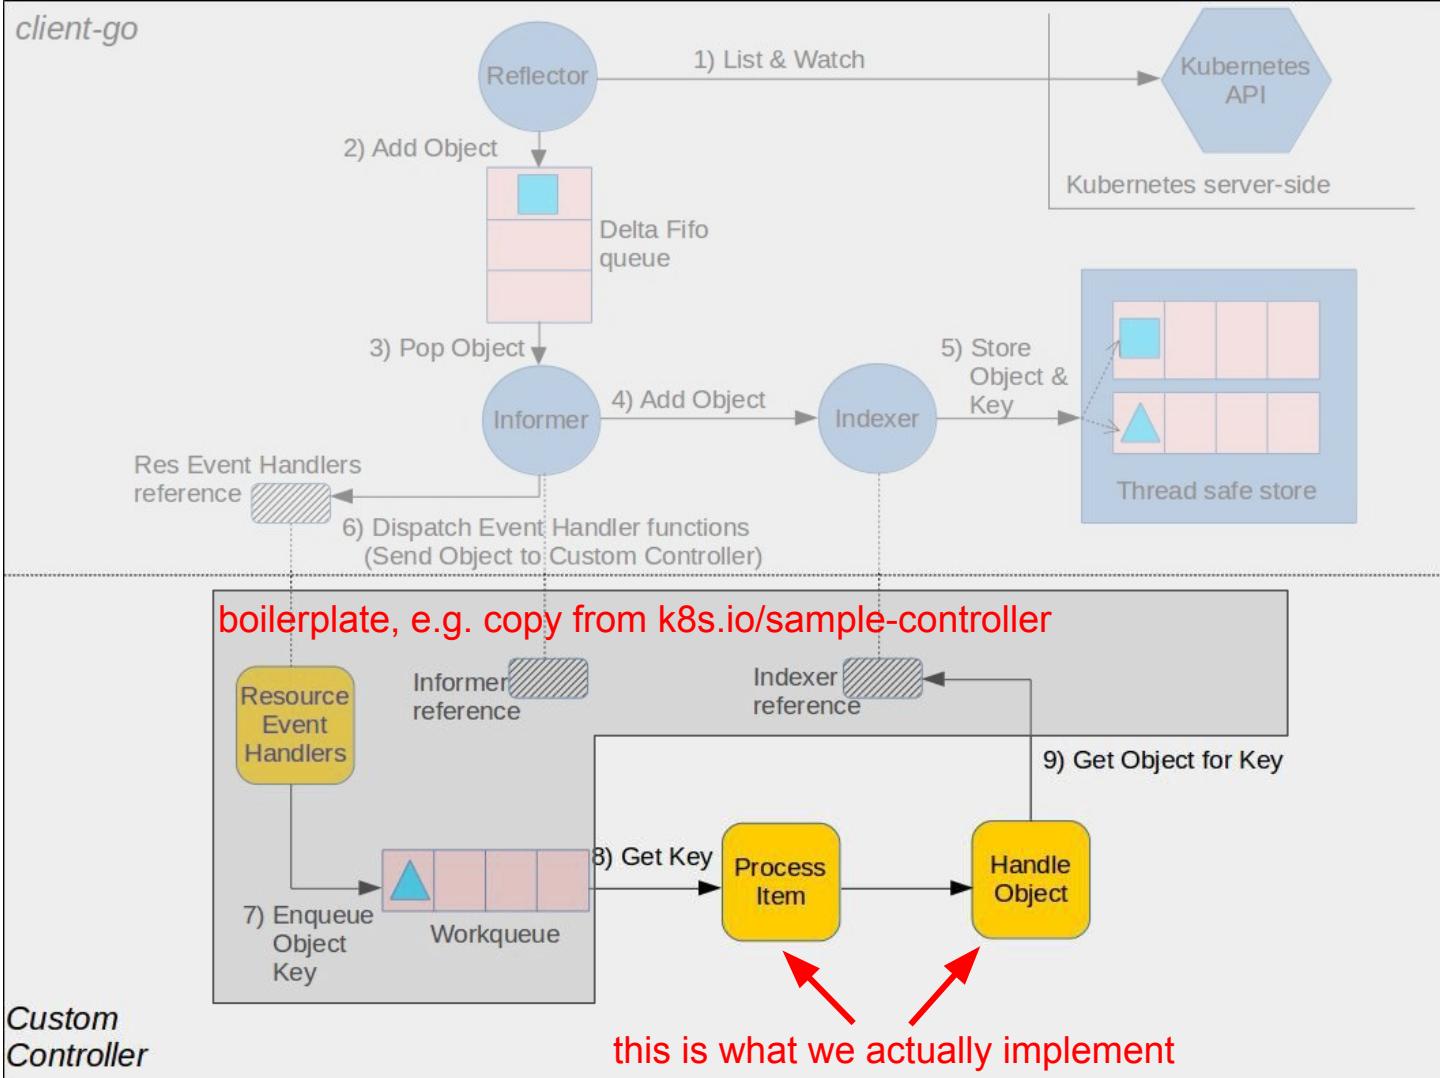


Image by @cloudark / Devdatta Kulkarni

<https://medium.com/@cloudark/kubernetes-custom-controllers-b6c7d0668fdf>

```
$ while true; do
    http GET http://127.0.0.1:8080/api/v1/namespaces/default |
        {jq ".metadata.annotations[\"example\"] = \"$RANDOM\""; sleep 1;} |
    http --check-status PUT http://127.0.0.1:8080/api/v1/namespaces/default || break
done
```

HTTP/1.1 409 Conflict

Content-Length: 310
Content-Type: application/json
Date: Fri, 10 Mar 2017 08:27:58 GMT
{

```
"apiVersion": "v1",
"code": 409,
"details": {
    "kind": "namespaces",
    "name": "default"
},
"kind": "Status",
"message": "Operation cannot be fulfilled on namespaces \"default\": the object has been
modified; please apply your changes to the latest version and try again",
"metadata": {},
"reason": "Conflict",
"status": "Failure"
}
```

„optimistic concurrency“

1. **read** object (preferably event driven with **watches**)
 2. **change** object, update the world
 3. **update** object on apiserver
 4. **repeat**
- expect version conflicts,
remember: **optimistic concurrency**

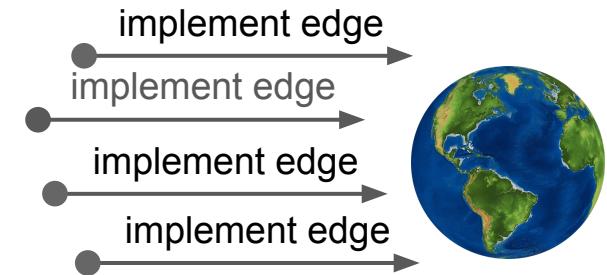


Level vs. Edge driven

1

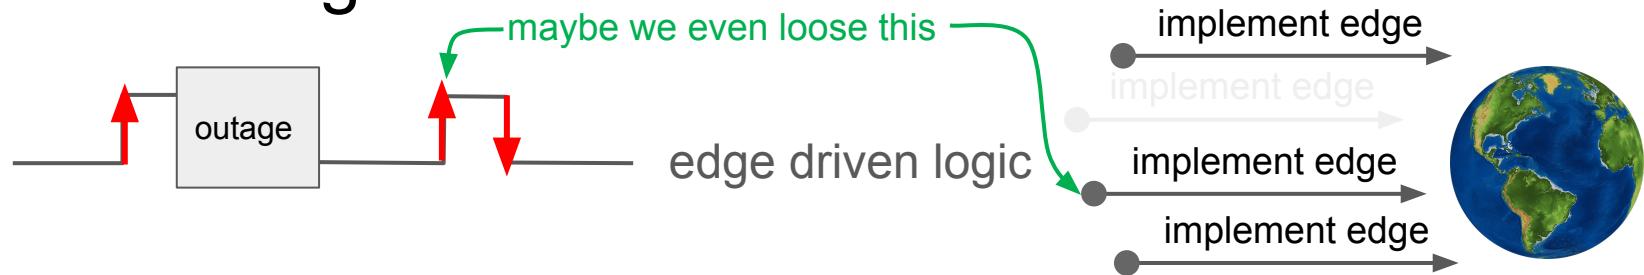


edge driven logic



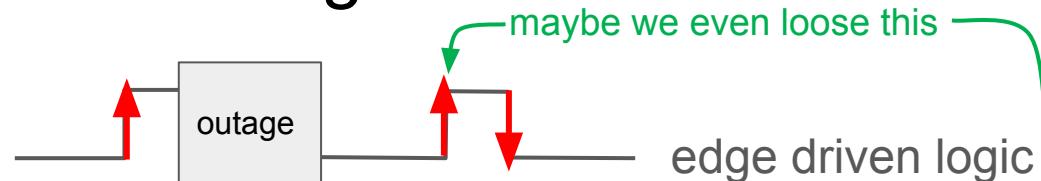
Level vs. Edge driven

1

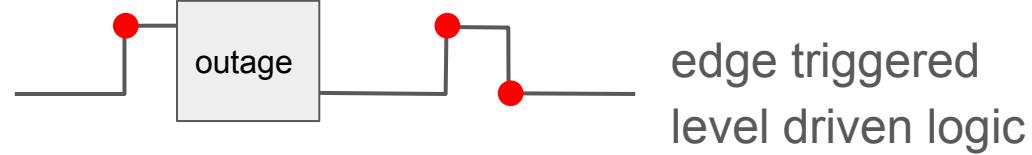


Level vs. Edge driven

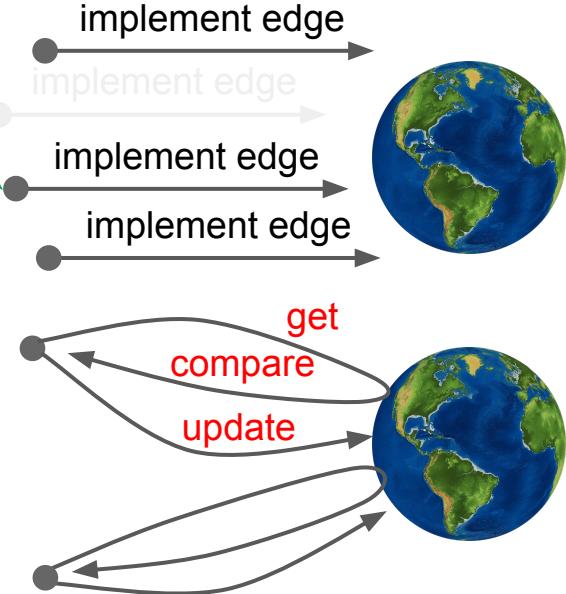
1



2



maybe we even loose this
edge driven logic



Level vs. Edge driven

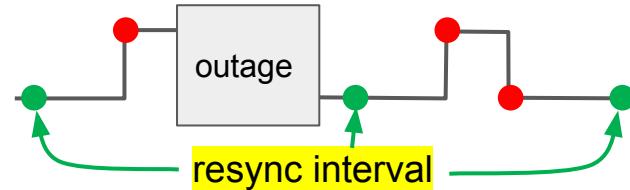
1



2



3

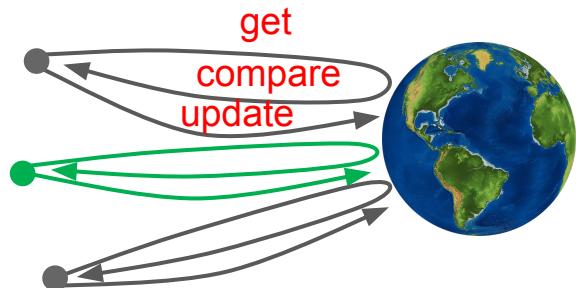
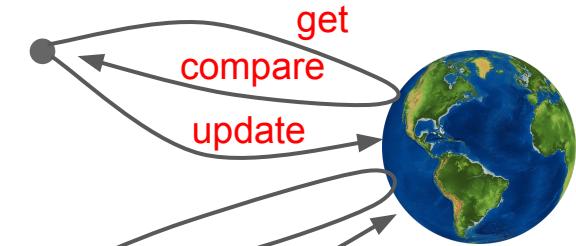
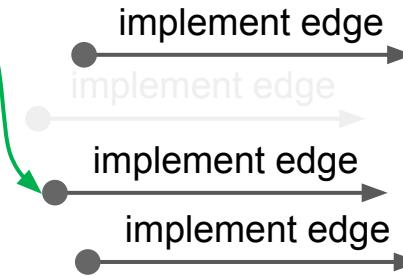


maybe we even loose this

edge driven logic

edge triggered
level driven logic

reconciliation
with resync



```
func (c *Controller) syncHandler(key string) error {
    // Convert the namespace/name string into a distinct namespace and name
    namespace, name, err := cache.SplitMetaNamespaceKey(key)
    if err != nil {
        runtime.HandleError(fmt.Errorf("invalid resource key: %s", key))
        return nil
    }

    // Get the health check policy resource with this namespace/name
    policy, err := c.healthCheckPolicyLister.HealthCheckPolicies(namespace).Get(name)
    if err != nil {
        if errors.NotFound(err) {
            runtime.HandleError(fmt.Errorf("health check policy '%s' in work queue no longer exists", key))
            return nil
        }
        return err
    }

    // keep original, and deep copy policy before mutating it
    original := policy
    policy = policy.DeepCopy()

    // ##### # TODO: add controller logic here #####
    policy.Status.PodsFailed = 42

    // skip update if there was no change
    if reflect.DeepEqual(original, policy) {
        return nil
    }

    // there was a change. Update resource on the server.
    if _, err := c.policyClient.PolicyV1alpha1().HealthCheckPolicies(policy.Namespace).Update(policy); errors.IsConflict(err) {
        c.workqueue.Add(key)
        return err
    } else if err != nil {
        return err
    }

    // no need to requeue, we will see our own update.

    return nil
}
```

split key into namespace+name

get object from cache

keep copy
deepcopy before mutation

update in memory

anything changed?

update on server

requeue on conflict

References for controllers

- k8s.io/sample-controller
- [Maciej Szulik's talk from KubeCon EU 2018](#)
- [Writing Kubernetes Custom Controllers](#) by Devdatta Kulkarni
- the controllers in core: k8s.io/kubernetes/pkg/controller



Katacoda

→ Writing controllers (step 6)

kubebuilder

```
kubebuilder init --domain k8s-go.openshift.org  
    --license apache2  
    --owner "The workshop members"
```

```
Run `dep ensure` to fetch dependencies (Recommended) [y/n]?  
y  
dep ensure  
Running make...  
make  
go generate ./pkg/... ./cmd/...  
go fmt ./pkg/... ./cmd/...  
go vet ./pkg/... ./cmd/...  
go run vendor/sigs.k8s.io/controller-tools/cmd/controller-gen/main.go all  
CRD manifests generated under '/Users/sts/Quellen/kubernetes/src/github.com/openshift/k8s-go/kubebuilder/config/crds'  
RBAC manifests generated under '/Users/sts/Quellen/kubernetes/src/github.com/openshift/k8s-go/kubebuilder/config/rbac'  
go test ./pkg/... ./cmd/... -coverprofile cover.out  
?      github.com/openshift/k8s-go/kubebuilder/pkg/apis      [no test files]  
?      github.com/openshift/k8s-go/kubebuilder/pkg/controller [no test files]  
?      github.com/openshift/k8s-go/kubebuilder/cmd/manager   [no test files]  
go build -o bin/manager github.com/openshift/k8s-go/kubebuilder/cmd/manager  
Next: Define a resource with:  
$ kubebuilder create api
```

```
kubebuilder create api --group policy*  
                      --version v1beta1  
                      --kind HealthCheckPolicy
```

```
Create Resource under pkg/apis [y/n]?  
y  
Create Controller under pkg/controller [y/n]?  
y  
Writing scaffold for you to edit...  
pkg/apis/policy/v1beta1/healthcheckpolicy_types.go  
pkg/apis/policy/v1beta1/healthcheckpolicy_types_test.go  
pkg/controller/healthcheckpolicy/healthcheckpolicy_controller.go  
pkg/controller/healthcheckpolicy/healthcheckpolicy_controller_test.go  
Running make...  
go generate ./pkg/... ./cmd/...
```

* needs some fixing of the controller code due to wrong import: <https://github.com/kubernetes-sigs/kubebuilder/issues/335>

```
import (
    kubescheme "k8s.io/client-go/kubernetes/scheme"
    kubebuilderclient "sigs.k8s.io/controller-runtime/pkg/client"
)

kubeconfig = flag.String("kubeconfig", "~/.kube/config", "path to the kubeconfig file")
flag.Parse()
config, err := clientcmd.BuildConfigFromFlags("", *kubeconfig)

client, err := kubebuilderclient.NewForConfig(config, Options{Scheme: kubescheme.Scheme})
pod := corev1.Pod{}
err := client.Get(context.TODO(), {"gophercon", "workshop"}, &pod)
```



Get, Create, Update, List, Delete, Patch, Watch



uses API discovery to know HTTP paths

```
func (r *ReconcileHealthCheckPolicy) Reconcile(request reconcile.Request) (reconcile.Result, error) {
    // Fetch the HealthCheckPolicy instance
    instance := &policyv1beta1.HealthCheckPolicy{}
    err := r.Get(context.TODO(), request.NamespacedName, instance)
    if err != nil {
        if errors.NotFound(err) {
            // Object not found, return. Created objects are automatically garbage collected.
            // For additional cleanup logic use finalizers.
            return reconcile.Result{}, nil
        }
        // Error reading the object - requeue the request.
        return reconcile.Result{}, err
    }
}
```

Namespace + Name

:

```
// TODO(user): Change this for the object type created by your controller
// Update the found object and write the result back if there are any changes
if !reflect.DeepEqual(deploy.Spec, found.Spec) {
    found.Spec = deploy.Spec
    log.Printf("Updating Deployment %s/%s\n", deploy.Namespace, deploy.Name)
    err = r.Update(context.TODO(), found)
    if err != nil {
        return reconcile.Result{}, err
    }
}
return reconcile.Result{}, nil
}
```

The Operator SDK

operator-sdk new app-operator

```
--api-version=policy.k8s-go.openshift.org/v1beta1  
--kind=HealthCheckPolicy
```

```
Create app-operator/.gitignore  
Create app-operator/cmd/app-operator/main.go  
Create app-operator/config/config.yaml  
Create app-operator/deploy/rbac.yaml  
Create app-operator/deploy/crd.yaml  
Create app-operator/deploy/cr.yaml  
Create app-operator/pkg/apis/policy/v1beta1/doc.go  
Create app-operator/pkg/apis/policy/v1beta1/register.go  
Create app-operator/pkg/apis/policy/v1beta1/types.go  
Create app-operator/pkg/stub/handler.go  
Create app-operator/tmp/build/build.sh  
Create app-operator/tmp/build/docker_build.sh  
Create app-operator/tmp/build/Dockerfile  
Create app-operator/tmp/codegen/boilerplate.go.txt  
Create app-operator/tmp/codegen/update-generated.sh  
Create app-operator/version/version.go  
Create app-operator/Gopkg.toml  
Run dep ensure ...
```

operator-sdk – workflow

```
$ operator-sdk build <docker-image-name>
```

```
$ docker push <docker-image-name>
```

```
$ OPERATOR_NAME=our-operator operator-sdk up local
```

```
INFO[0000] Go Version: go1.10.3
```

```
INFO[0000] Go OS/Arch: darwin/amd64
```

```
INFO[0000] operator-sdk Version: 0.0.5+git
```

```
INFO[0003] Metrics service foo created
```

```
INFO[0003] Watching policy.k8s-go.openshift.org/v1beta1, HealthCheckPolicy, default, 5
```

```
import (
    "github.com/operator-framework/operator-sdk/pkg/sdk"
    corev1 "k8s.io/api/core/v1"
)
```

```
pod := corev1.Pod{
    TypeMeta: {ApiVersion: "v1", Kind: "Pod"},
    ObjectMeta: {Name: "busybox", Namespace: "default"},
    Spec: ...
}
err := sdk.Create(&pod)
```

Get, Create, Update, List, Delete, Patch, Watch

must be filled in

uses API discovery to know HTTP paths

```
func NewHandler() sdk.Handler {
    return &Handler{}
}

type Handler struct {
    // Fill me
}

func (h *Handler) Handle(ctx context.Context, event sdk.Event) error {
    switch o := event.Object.(type) {
    case *v1beta1.HealthCheckPolicy:
        err := sdk.Create(newbusyBoxPod(o)) ← Client call
        if err != nil && !errors.Exists(err) {
            logrus.Errorf("Failed to create busybox pod : %v", err)
            return err
        }
    }
    return nil
}
```

```
type Event struct {
    Object  Object
    Deleted bool
}
```

References for operator-sdk

<https://github.com/operator-framework/operator-sdk>

[A complete guide to Kubernetes Operator SDK](#) by Toader Sebastian

[Building an Kubernetes Operator for Prometheus and Thanos](#) by Rob Szumski

Agenda

- **API basics:** apigroups, resource vs. kind, versioning, discovery, **curl scenario**
- **Client-go basics:** clientsets, kubeconfig, **client-go scenario**, informers, **informer scenario**
- **Kubernetes objects in Golang:** ObjectMeta, TypeMeta
- **API Machinery:** Scheme, GVK, GVR, RestMapper
- **CRDs from user point of view:** CRD yaml, Validation, **CRD yaml scenario**
- **Code-Generation:** deepcopy-gen, client-gen, informer-gen, lister-gen
- **CRDs from developer point of view:** types.go, register.go, pkg/apis, code-generator, **CRD Go scenario**
- **Controllers:** controller loop, optimistic concurrency, edge-vs-level driven, **controller scenario**
- **Kubebuilder:** init, create api, **kubebuilder scenario – trainer led**
- **Operator SDK:** new operator, **operator scenario – trainer led**
- **Other topics:** deployment / in-cluster config, RBAC rules, service account, subresources, versioning, **misc scenario ("deploy your controller in the cluster")**

Quick glimpse on other Topics

Deployment vs. in-cluster config

```
import (
    "k8s.io/client-go/tools/clientcmd"
)

kubeconfig = flag.String("kubeconfig", "~/.kube/config", "path to the kubeconfig file")
flag.Parse()

config, err := rest.InClusterConfig()
if err != nil {
    config, err := clientcmd.BuildConfigFromFlags("", *kubeconfig)
    if err != nil { ... }
}

clientset, err := kubernetes.NewForConfig(config)
```

looks at /var/run/secrets/kubernetes.io/serviceaccount/token
and consider:

- \$KUBERNETES_SERVICE_HOST
- \$KUBERNETES_SERVICE_PORT

client-go dynamic client – unstructured.Unstructured

```
import (
    metav1 "k8s.io/apimachinery/pkg/apis/meta/v1"
    "k8s.io/client-go/tools/clientcmd"
    "k8s.io/client-go/dynamic"
)
```

```
kubeconfig = flag.String("kubeconfig", "~/.kube/config", "path to the kubeconfig file")
flag.Parse()
config, err := clientcmd.BuildConfigFromFlags("", *kubeconfig)
client, err := dynamic.NewForConfig(config)

gvr := runtime.GroupVersionResource{Group: "", Version: "v1", Resource: "pods"}
pod, err := client.Resource(gvr).Namespace("gophercon").Get("workshop", {})
```

↑
*unstructured.Unstructured

↑
Get, Create, Update, List, Delete, Patch, Watch

RBAC rules

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: policy-controller-role
rules:
- apiGroups:           - apiGroups:
  - apps               - policy.k8s-go.openshift.org
resources:          resources:
  - deployments        - healthcheckpolicies
verbs:              verbs:
  - get                - get
  - list               - list
  - watch              - watch
  - create             - create
  - update             - update
  - patch              - patch
  - delete             - delete
```

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: policy-controller-binding
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: policy-controller-role
subjects:
- namespace: controller-namespace
  kind: ServiceAccount
  name: default
```

every namespace has a ServiceAccount named default

User or Group or ServiceAccount

Custom Service Account

```
$ kubectl create serviceaccount policy-controller
```

```
apiVersion: v1
kind: Pod
metadata:
  name: policy-controller
spec:
  serviceAccountName: policy-controller
....
```

Subresources – /scale

```
$ kubectl scale --replicas=3 <your-custom-resource> --v=7  
I0429 21:17:53.138353 66743 round_tripper.go:383] PUT
```

https://<host>/apis/<group>/v1/<your-custom-resource>/scale

apiVersion: apiextensions.k8s.io/v1beta1

kind: CustomResourceDefinition

spec:

subresources:

scale:

specReplicasPath: .spec.replicas

statusReplicasPath: .status.replicas

labelSelectorPath: .status.labelSelector

Subresources – /status

- **separate RBAC rules** for controller and users
 - controller writes to /apis/<group>/v1/<resource>/status
 - user writes to /apis/<group>/v1/<resource>
- ⇒ **The user cannot mutate the status.**
- on /status all fields outside of .status are ignored
- on / all fields under .status are ignored

apiVersion: apiextensions.k8s.io/v1beta1

kind: CustomResourceDefinition

spec:

subresources:

status: {} ← enabled like this, available in Kubernetes 1.11 as beta

Kubernetes Style Versioning

v1alpha1

v1alpha2

v1beta1

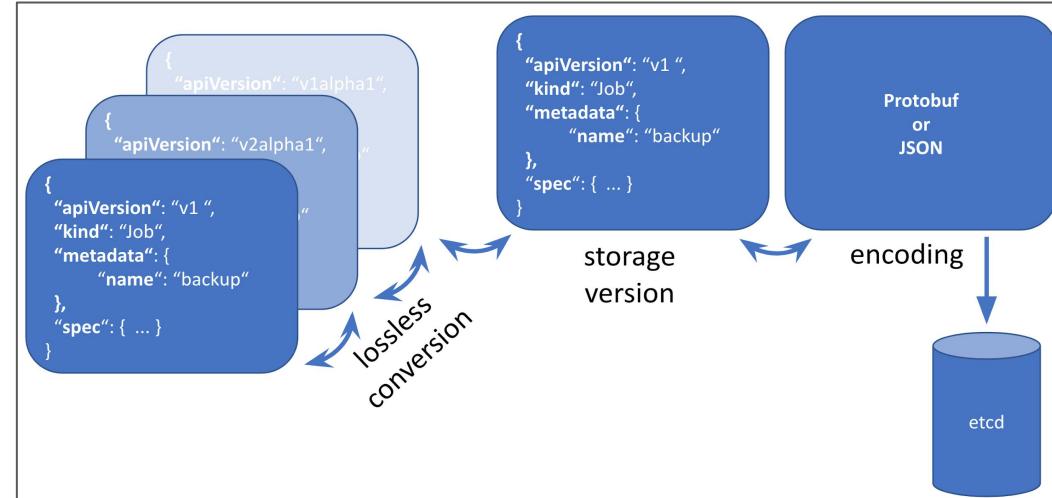
v1

v2alpha1

v2beta1

v2

- v1alpha1 – unstable:
 - might go away or change any time
 - often disabled by default
- v1beta1 – towards stable:
 - exists at least one release in parallel to v1
 - no incompatible API change
- v1 – stable, GA:
 - will stay
 - will be compatible



CRD Versions in Kubernetes 1.11 and beyond

`apiVersion: apiextensions.k8s.io/v1beta1`

`kind: CustomResourceDefinition`

`spec:`

`version: v1alpha1` ← will be slowly deprecated

`versions:` ← in Kubernetes 1.11

- `name: v1alpha1`

`served: true`

`storage: false`

- `name: v1beta1`

`served: true`

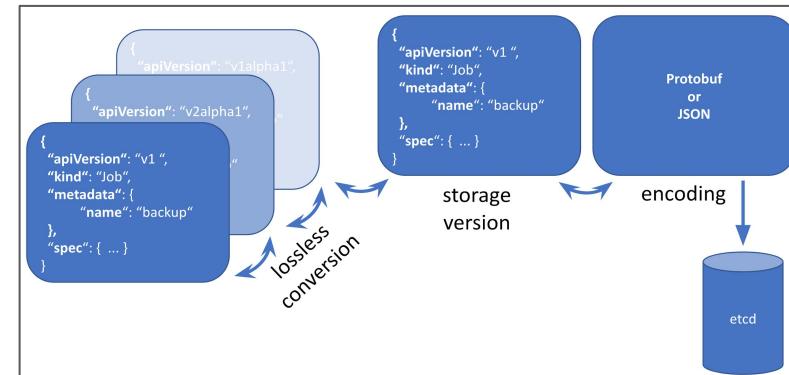
`storage: true`

`conversion:`

`webhook:` ← planned in Kubernetes 1.12 as alpha

`clientConfig:`

`url: ...`



Agenda

- **API basics:** apigroups, resource vs. kind, versioning, discovery, **curl scenario**
- **Client-go basics:** clientsets, kubeconfig, **client-go scenario**, informers, **informer scenario**
- **Kubernetes objects in Golang:** ObjectMeta, TypeMeta
- **API Machinery:** Scheme, GVK, GVR, RestMapper
- **CRDs from user point of view:** CRD yaml, Validation, **CRD yaml scenario**
- **Code-Generation:** deepcopy-gen, client-gen, informer-gen, lister-gen
- **CRDs from developer point of view:** types.go, register.go, pkg/apis, code-generator, **CRD Go scenario**
- **Controllers:** controller loop, optimistic concurrency, edge-vs-level driven, **controller scenario**
- **Kubebuilder:** init, create api, **kubebuilder scenario – trainer led**
- **Operator SDK:** new operator, **operator scenario – trainer led**
- **Other topics:** deployment / in-cluster config, RBAC rules, service account, subresources, versioning, **misc scenario ("deploy your controller in the cluster")**

Backup

Zoom into apiextensions-apiserver

